



Image-Based Situation Awareness Audit 8.5.2018

Sakari Lampola



Previous Audit 28.2.2018

Next steps

- Kalman filter parameter adjustments (Q1) ●
- Dataset selection (Q1) ●
- Stereo vision (Q2) ●
- Camera yaw, pitch, roll estimation (Q2) ●
- Speech recognition (Q2) ○
- Semantic segmentation (Q2) ○
- Experiments in the wild (Q2)
- Paper (Q3) ●
- Speech analysis (Q3) ●
- Speech generation (Q3) ●
- Use cases (Q4) ●



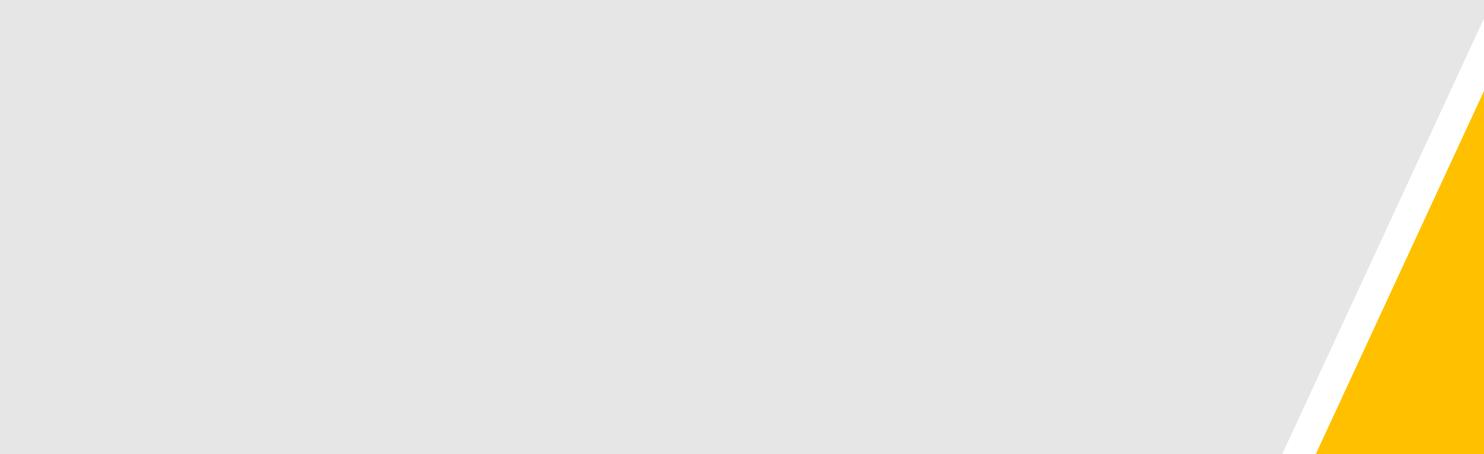
Other

- Body forecast
 - kinetic
 - based on class history
 - based on swarm history
- R matrix estimation
- Monograph or papers?

Project Plan

	2018	2019	2020	2021
Methodology				
Preparation of research infra				
Method survey				
Building test cases				
Testing and comparison				
Prototype				
Definition				
Planning				
Implementation				
Testing and fixing				
Method follow-up				
Writing thesis				
Dissertation				

1. Methodology / Preparation of research infra
 - a. Software platforms are constructed and tested
 - b. Off-the-shelf models are acquired and tested
 - c. Necessary skills on platforms are learned
2. Methodology / Method survey
 - a. Current state-of-art methods are studied
 - b. Methods are constructed and tested on the software platforms
3. Method follow-up
 - a. Screening of conference papers related to the subject
 - b. Possibly integrating new methods to the project



Work Done

Dataset Selection

Specification:

- Video
- Stereo
- Distance information
- Outdoor + indoor
- Odometry

Select category: [City](#) | [Residential](#) | [Road](#) | [Campus](#) | [Person](#) | [Calibration](#)

Data Category: City

Before browsing, please wait some moments until this page is fully loaded.



2011_09_26_drive_0001 (0.4 GB)

Length: 11 frames (00:01 minutes)

Image resolution: 1392 x 512 pixels

Labels: 12 Cars, 0 Vans, 0 Trucks, 0 Pedestrians, 0 Sitters, 2 Cyclists, 1 Trams, 0 Misc

Downloads: [unsynced+unrectified data] [synced+rectified data] [calibration] [tracklets]



2011_09_26_drive_0002 (0.3 GB)

Length: 83 frames (00:08 minutes)

Image resolution: 1392 x 512 pixels

Labels: 1 Cars, 0 Vans, 0 Trucks, 0 Pedestrians, 0 Sitters, 2 Cyclists, 0 Trams, 0 Misc

Downloads: [unsynced+unrectified data] [synced+rectified data] [calibration] [tracklets]



2011_09_26_drive_0005 (0.6 GB)

Length: 160 frames (00:16 minutes)

Image resolution: 1392 x 512 pixels

Labels: 9 Cars, 3 Vans, 0 Trucks, 2 Pedestrians, 0 Sitters, 1 Cyclists, 0 Trams, 0 Misc

Downloads: [unsynced+unrectified data] [synced+rectified data] [calibration] [tracklets]

The KITTI Vision Benchmark Suite

A project of Karlsruhe Institute of Technology
and Toyota Technological Institute at Chicago



[home](#) [setup](#) [stereo](#) [flow](#) [sceneflow](#) [depth](#) [odometry](#) [object](#) [tracking](#) [road](#) [semantics](#) [raw data](#) [submit results](#)

Andreas Geiger (MPI Tübingen) | Philip Lenz (KIT) | Christoph Stiller (KIT) | Raquel Urtasun (University of Toronto)

Raw Data

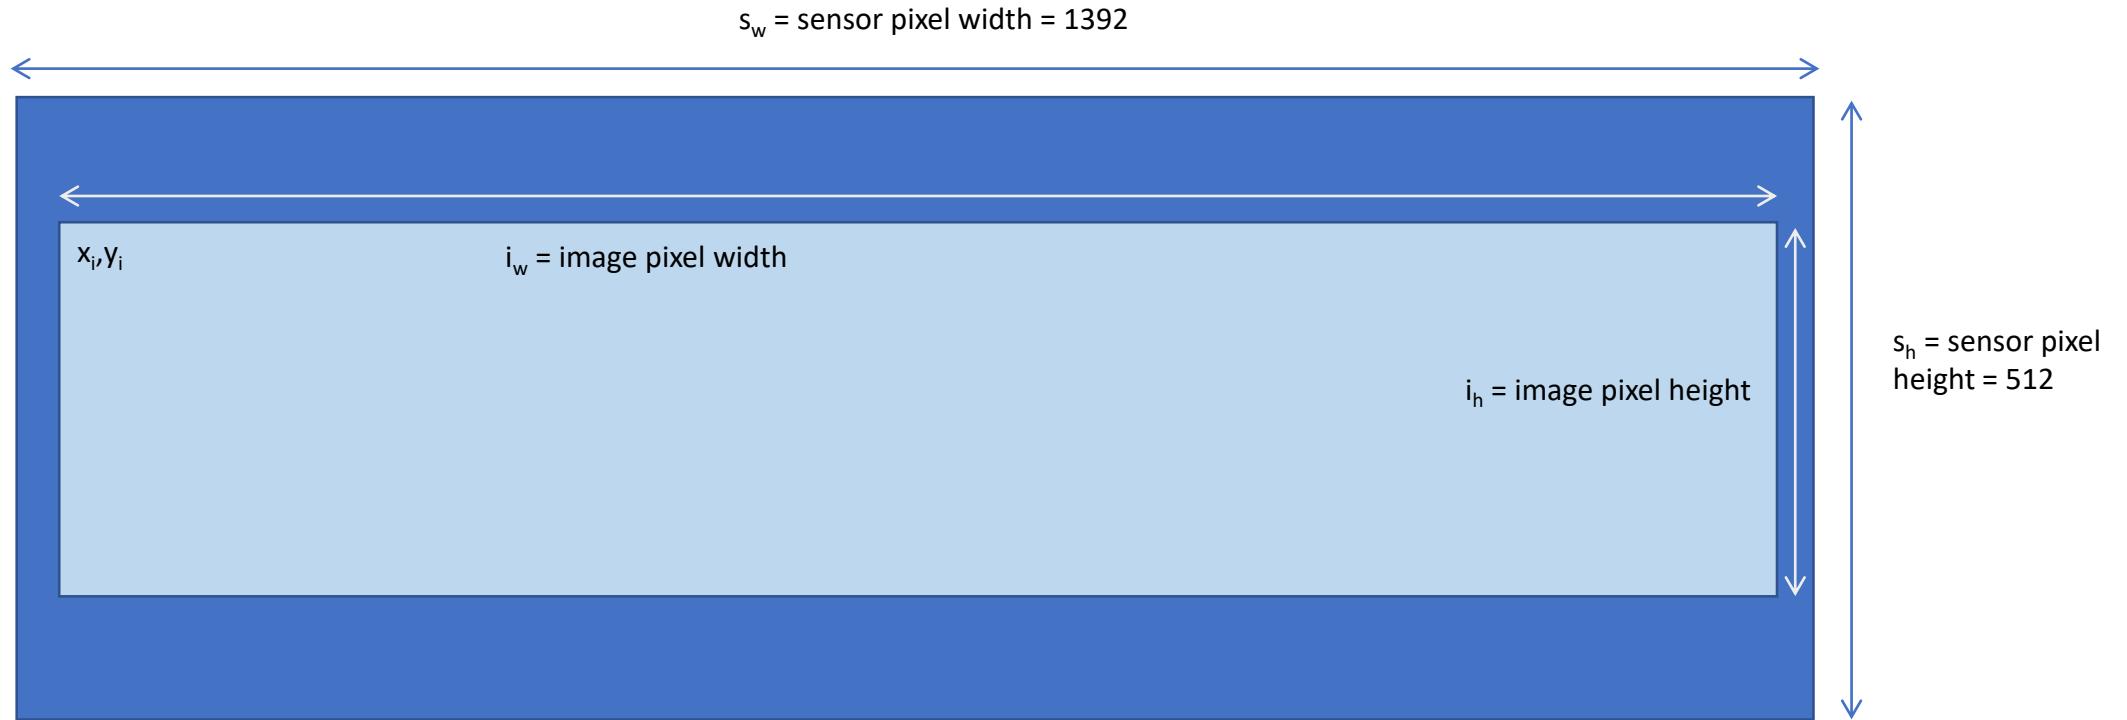
This page contains our raw data recordings, sorted by category (see menu above). So far, we included only sequences, for which we either have 3D object labels or which occur in our odometry benchmark training set. The dataset comprises the following information, captured and synchronized at 10 Hz:

- Raw (unsynced+unrectified) and processed (synced+rectified) grayscale stereo sequences (0.5 Megapixels, stored in png format)
- Raw (unsynced+unrectified) and processed (synced+rectified) color stereo sequences (0.5 Megapixels, stored in png format)
- 3D Velodyne point clouds (100k points per frame, stored as binary float matrix)
- 3D GPS/IMU data (location, speed, acceleration, meta information, stored as text file)
- Calibration (Camera, Camera-to-GPS/IMU, Camera-to-Velodyne, stored as text file)
- 3D object tracklet labels (cars, trucks, trams, pedestrians, cyclists, stored as xml file)



Open question: Indoor? Self generated?

Image format / cropping



Assumptions:

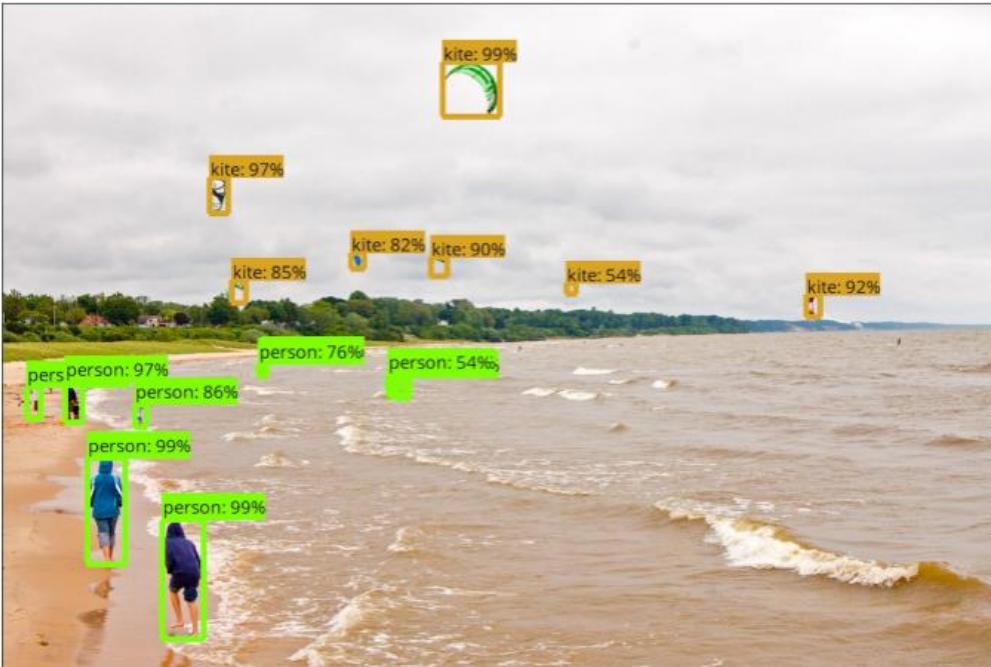
- image is symmetrically cropped (optical center fixed)
- rectification ignored

$$x_i = (s_w - i_w)/2$$

$$y_i = (s_h - i_h)/2$$

Tensorflow Object Detection API

Creating accurate machine learning models capable of localizing and identifying multiple objects in a single image remains a core challenge in computer vision. The TensorFlow Object Detection API is an open source framework built on top of TensorFlow that makes it easy to construct, train and deploy object detection models. At Google we've certainly found this codebase to be useful for our computer vision needs, and we hope that you will as well.



COCO-trained models {#coco-models}

Model name	Speed (ms)	COCO mAP[^1]	Outputs
ssd_mobilenet_v1_coco	30	21	Boxes
ssd_inception_v2_coco	42	24	Boxes
faster_rcnn_inception_v2_coco	58	28	Boxes
faster_rcnn_resnet50_coco	89	30	Boxes
faster_rcnn_resnet50_lowproposals_coco	64		Boxes
rfcn_resnet101_coco	92	30	Boxes
faster_rcnn_resnet101_coco	106	32	Boxes
faster_rcnn_resnet101_lowproposals_coco	82		Boxes
faster_rcnn_inception_resnet_v2_atrous_coco	620	37	Boxes
faster_rcnn_inception_resnet_v2_atrous_lowproposals_coco	241		Boxes
faster_rcnn_nas	1833	43	Boxes
faster_rcnn_nas_lowproposals_coco	540		Boxes
mask_rcnn_inception_resnet_v2_atrous_coco	771	36	Masks
mask_rcnn_inception_v2_coco	79	25	Masks
mask_rcnn_resnet101_atrous_coco	470	33	Masks
mask_rcnn_resnet50_atrous_coco	343	29	Masks

One of these will be the final model.

Kitti-trained models {#kitti-models}

Model name	Speed (ms)	Pascal mAP@0.5 (ms)	Outputs
faster_rcnn_resnet101_kitti	79	87	Boxes

Lottery prize!!!! Will be used to implement localization and velocity estimation

Open Images-trained models {#open-images-models}

Model name	Speed (ms)	Open Images mAP@0.5[^2]	Outputs
faster_rcnn_inception_resnet_v2_atrous_oid	727	37	Boxes
faster_rcnn_inception_resnet_v2_atrous_lowproposals_oid	347		Boxes

News

- 2017 Challenge Winners for Detection, Keypoint, & Stuff tasks have been announced! Please visit the [Joint COCO and Places Recognition ICCV workshop page](#) for details.
- This website is now hosted on [Github](#), which provides page source and history.
- Keypoint analysis tools are now available, see [keypoints evaluation](#), Section 4.

What is COCO?



COCO is a large-scale object detection, segmentation, and captioning dataset. COCO has several features:

- ✓ Object segmentation
- ✓ Recognition in context
- ✓ Superpixel stuff segmentation
- ✓ 330K images (>200K labeled)
- ✓ 1.5 million object instances
- ✓ 80 object categories
- ✓ 91 stuff categories
- ✓ 5 captions per image
- ✓ 250,000 people with keypoints

Research Paper

Download the paper that describes the Microsoft COCO dataset.

Collaborators

Tsung-Yi Lin Google Brain

Genevieve Patterson MSR

Matteo R. Ronchi Caltech

Yin Cui Cornell Tech

Michael Maire TTI-Chicago

Serge Belongie Cornell Tech

Lubomir Bourdev WaveOne, Inc.

Ross Girshick FAIR

James Hays Georgia Tech

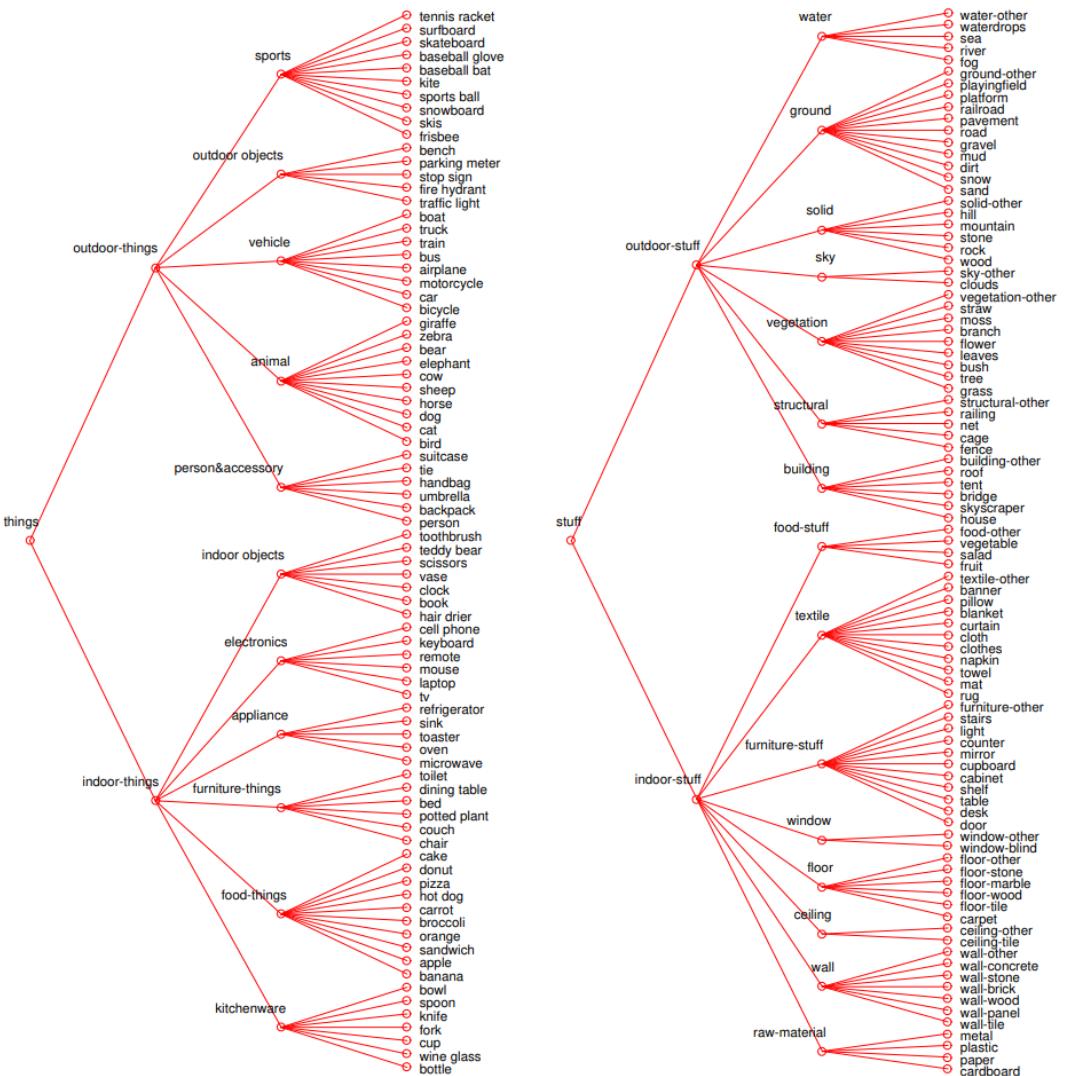
Pietro Perona Caltech

Deva Ramanan CMU

Larry Zitnick FAIR

Piotr Dollár FAIR

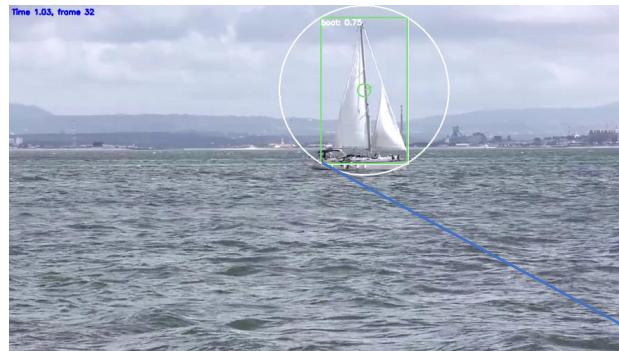
Sponsors



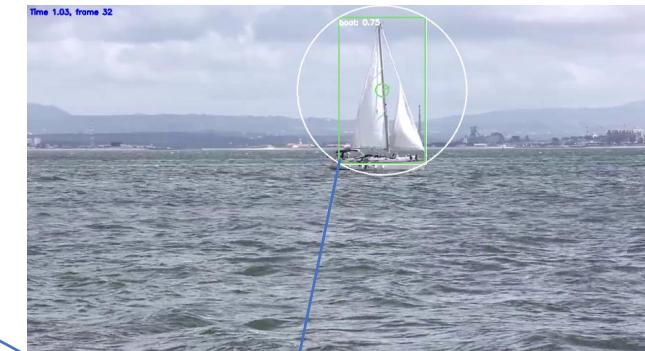
 Download
paper here

Stereo Vision

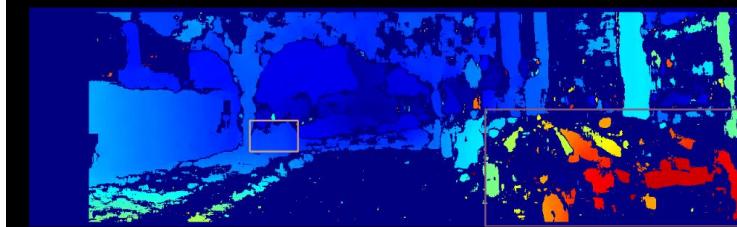
Left



Right



disparity d





ZED™

2K Stereo Camera

The World's First
3D Camera for Depth Sensing
and Motion Tracking

[ORDER FOR \\$449](#)

Shipping Worldwide

Compatible OS



Windows 7, 8, 10



Linux

Third-party Support



SDK System Requirements

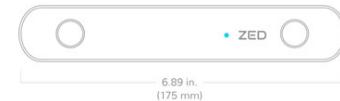
- Dual-core 2.3GHz or faster processor
- 4 GB RAM or more
- Nvidia GPU with compute capability > 3.0

In The Box

- ZED Stereo camera
- Mini Tripod stand
- USB Drive with Drivers and SDK
- Documentation



Dimensions



Features

- High-Resolution and High Frame-rate 3D Video Capture
- Depth Perception indoors and outdoors at up to 20m
- 6-DoF Positional Tracking
- Spatial Mapping

Video

Video Mode	Frames per second	Output Resolution (side by side)
2.2K	15	4416x1242
1080p	30	3840x1080
720p	60	2560x720
WVGA	100	1344x376

Depth

Depth Resolution
Same as selected video resolution

Depth Format
32-bits


Depth Range
0.5 - 20 m (2.3 to 65 ft)

Stereo Baseline
120 mm (4.7")

Mapping left and right image patterns



```
class,confidence,x,y,width,height,hue0,hue1,hue2,saturation,value  
-----  
2,1.00,566.00,215.50,52.00,137.00,0.315,0.437,0.247,71.996,124.843  
-----  
2,1.00,530.50,216.00,49.00,134.00,0.291,0.468,0.242,72.229,121.822
```

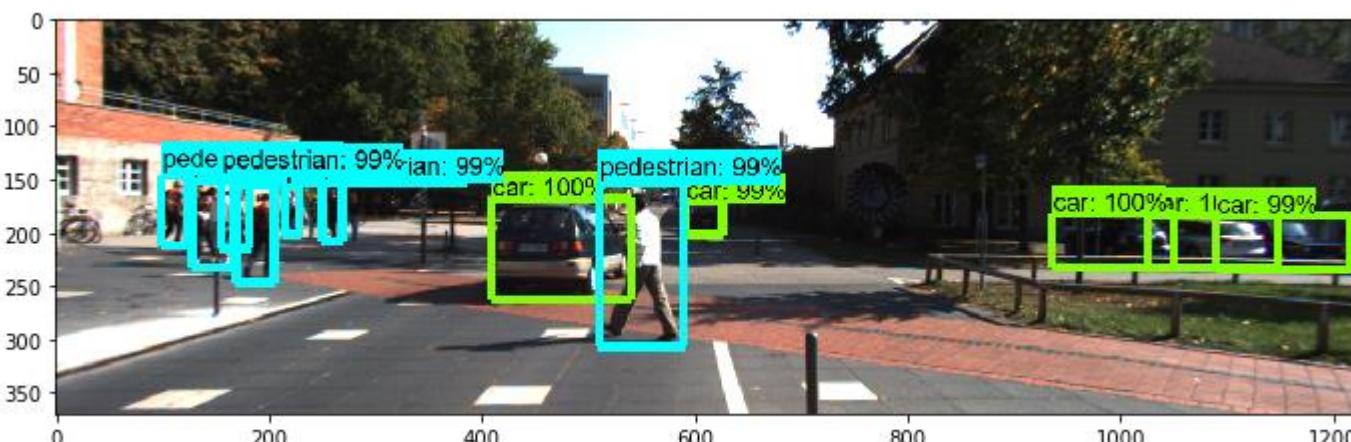
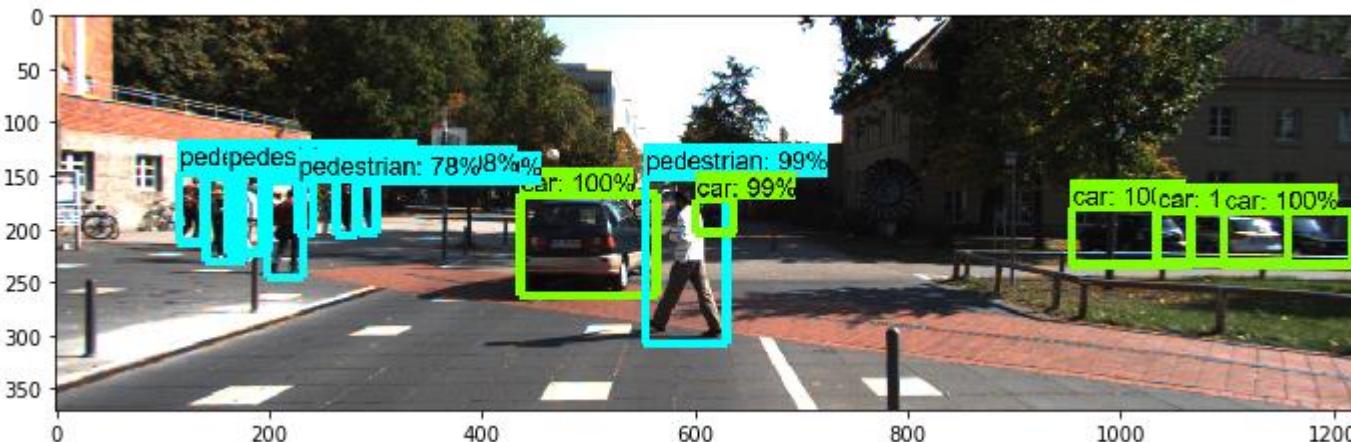
Pattern features to match



x,y = bounding box center location
saturation,value = mean values

3-bin hue histogram:





```

class,confidence,x,y,width,height,hue0,hue1,hue2,saturation,value
-----
1,1.00,1010.50,206.50,113.00,53.00,0.354,0.438,0.208,74.793,48.389
1,1.00,1096.00,207.00,124.00,48.00,0.313,0.514,0.173,75.248,79.118
1,1.00,1157.50,208.50,121.00,49.00,0.274,0.551,0.175,79.360,81.186
1,1.00,500.50,215.00,129.00,94.00,0.367,0.389,0.244,75.836,70.045
2,1.00,592.00,229.50,78.00,155.00,0.550,0.254,0.196,59.824,104.288
1,1.00,619.50,187.50,37.00,31.00,0.286,0.323,0.391,58.958,58.623
2,1.00,159.00,190.50,36.00,79.00,0.434,0.220,0.346,77.767,116.533
2,1.00,293.50,180.50,17.00,45.00,0.354,0.299,0.346,78.097,70.915
2,1.00,127.50,180.00,25.00,64.00,0.493,0.174,0.333,97.595,144.548
2,1.00,216.00,200.00,34.00,90.00,0.422,0.256,0.322,82.717,112.057
2,1.00,184.50,184.00,33.00,72.00,0.413,0.229,0.358,68.688,159.507
2,1.00,177.50,185.50,33.00,75.00,0.395,0.244,0.361,70.105,148.906
2,0.98,272.50,179.50,19.00,51.00,0.397,0.261,0.342,80.279,69.196
2,0.79,235.00,180.50,16.00,47.00,0.483,0.209,0.309,94.992,127.202
-----
1,1.00,1087.00,207.50,124.00,49.00,0.292,0.507,0.201,71.784,76.482
1,1.00,993.50,207.00,115.00,50.00,0.348,0.410,0.242,72.279,47.679
1,1.00,475.00,215.00,132.00,94.00,0.311,0.366,0.323,69.133,74.729
1,1.00,609.00,186.00,36.00,30.00,0.144,0.298,0.557,59.536,30.091
2,1.00,550.50,230.00,81.00,152.00,0.564,0.262,0.175,60.986,116.519
2,1.00,187.50,200.50,37.00,91.00,0.394,0.285,0.322,78.999,107.080
2,1.00,147.50,191.50,39.00,81.00,0.381,0.260,0.359,69.986,128.833
2,1.00,222.50,178.50,15.00,45.00,0.376,0.330,0.293,78.055,120.033
2,1.00,260.00,180.50,20.00,55.00,0.394,0.267,0.339,79.143,69.303
2,1.00,112.00,178.50,26.00,65.00,0.509,0.172,0.319,90.011,143.473
1,1.00,1153.50,209.00,127.00,50.00,0.209,0.593,0.198,77.370,78.260
2,0.99,169.00,180.50,28.00,67.00,0.432,0.317,0.251,75.630,159.326

```

Pattern matching based on feature difference is required!

Probabilistic model, answering the question: What is the probability two patterns represent the same object?

Feature vector F:

- confidence
- x
- y
- width
- height
- hue0
- hue1
- hue2
- saturation
- value

Assumption:

$$P(i \text{ and } j \text{ are same pattern}) \sim N(F_i - F_j | \mu_F, V_F)$$

Note: Class is not included as
it is **required** to be the same

μ_F , V_F were estimated by matching 84 patterns in 28 KITTI stereo image pairs representing city, residential, campus and person categories, including both cars and pedestrians.

```
In [25]: mean = df.mean()
```

```
In [26]: mean
```

```
Out[26]:
```

	dConfidence	-0.002024
dX	30.119048	
dY	0.142857	
dWidth	1.595238	
dHeight	1.071429	
dHue0	0.020762	
dHue1	-0.012524	
dHue2	-0.008333	
dSaturation	3.859619	
dValue	-1.029405	
dtype:	float64	

Note: Mean disparity (dX) is appr. 30 pixels

```
In [21]: covariance=df.cov()
```

```
In [22]: covariance
```

```
Out[22]:
```

	dConfidence	dX	dY	dWidth	dHeight	dHue0	dHue1	dHue2	dSaturation	dValue
dConfidence	0.002886	-0.069214	-0.022539	0.116641	-0.000818	-0.000207	0.000077	0.000131	0.046382	-0.003060
dX	-0.069214	411.545898	1.482788	53.121056	-1.707401	0.007306	-0.041094	0.034841	-31.935069	5.324796
dY	-0.022539	1.482788	6.991394	-20.158348	-2.624785	0.019775	-0.025219	0.005289	-2.481722	2.036318
dWidth	0.116641	53.121056	-20.158348	559.761905	75.860585	0.054818	-0.032215	-0.022245	4.598916	-36.784720
dHeight	-0.000818	-1.707401	-2.624785	75.860585	30.356282	0.030632	-0.030119	-0.000337	2.336329	-2.746706
dHue0	-0.000207	0.007306	0.019775	0.054818	0.030632	0.001820	-0.001190	-0.000622	-0.019008	0.125514
dHue1	0.000077	-0.041094	-0.025219	-0.032215	-0.030119	-0.001190	0.001925	-0.000741	0.040519	-0.041532
dHue2	0.000131	0.034841	0.005289	-0.022245	-0.000337	-0.000622	-0.000741	0.001361	-0.021563	-0.083944
dSaturation	0.046382	-31.935069	-2.481722	4.598916	2.336329	-0.019008	0.040519	-0.021563	24.449793	2.102903
dValue	-0.003060	5.324796	2.036318	-36.784720	-2.746706	0.125514	-0.041532	-0.083944	2.102903	78.502261

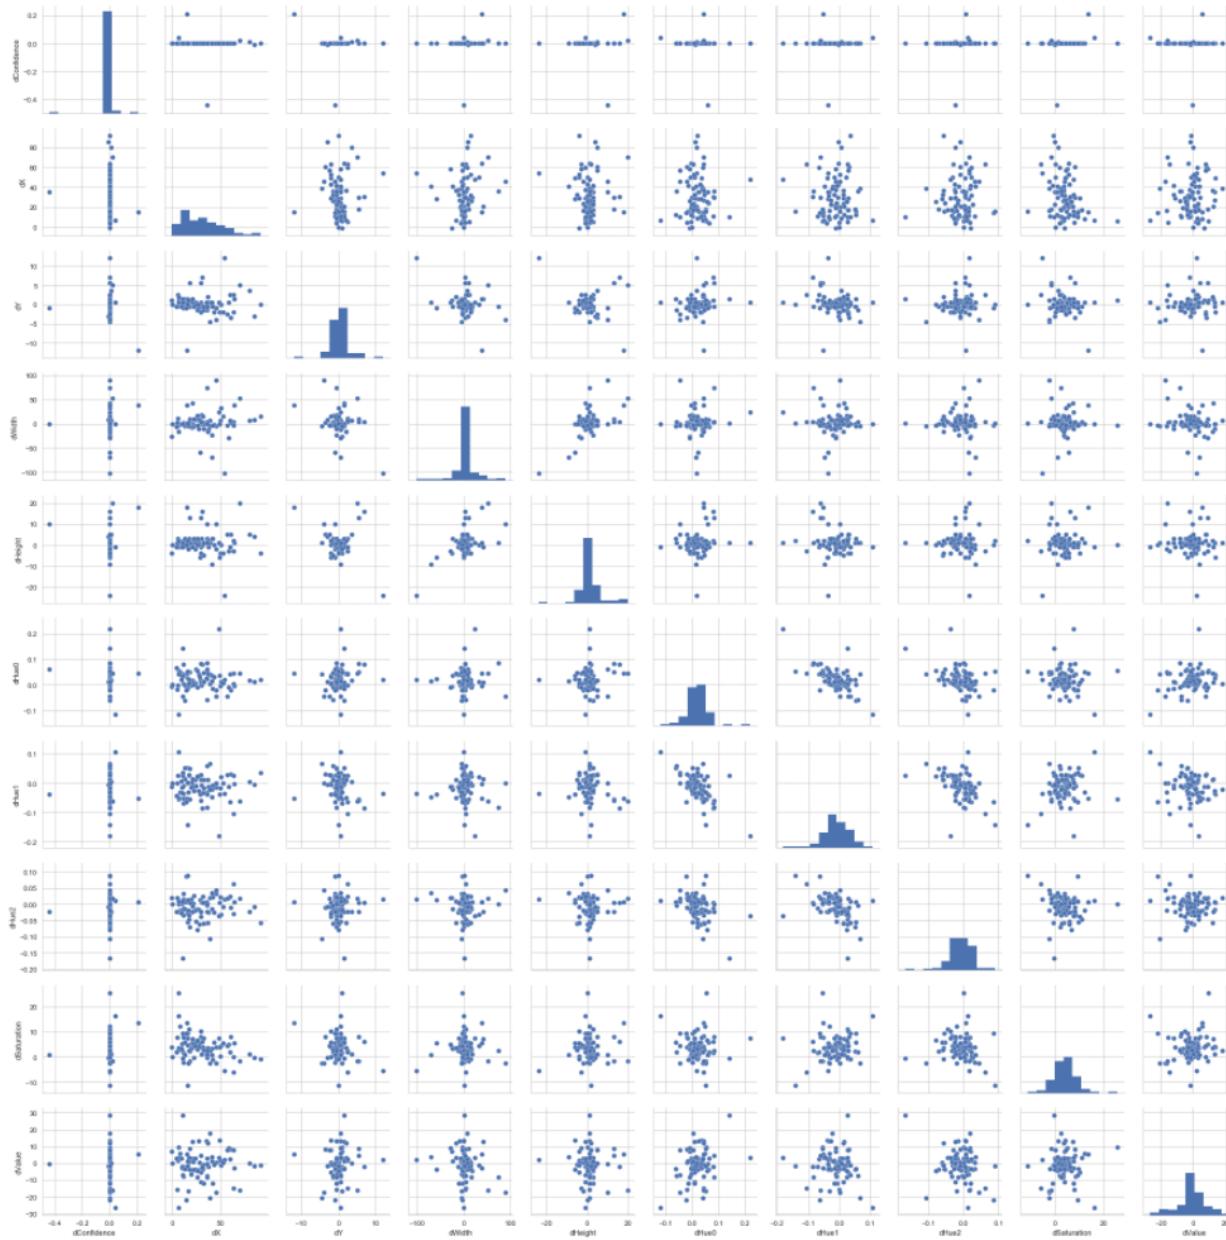
```
In [24]: df.describe()
```

Out[24]:

	dConfidence	dX	dY	dWidth	dHeight	dHue0	dHue1	dHue2	dSaturation	dValue
count	84.000000	84.000000	84.000000	84.000000	84.000000	84.000000	84.000000	84.000000	84.000000	84.000000
mean	-0.002024	30.119048	0.142857	1.595238	1.071429	0.020762	-0.012524	-0.008333	3.859619	-1.029405
std	0.053724	20.286594	2.644124	23.659288	5.509654	0.042662	0.043879	0.036886	4.944673	8.860150
min	-0.440000	-1.000000	-12.000000	-102.000000	-24.000000	-0.117000	-0.181000	-0.166000	-11.327000	-26.690000
25%	0.000000	14.875000	-0.500000	-3.250000	-1.000000	-0.001250	-0.035250	-0.024500	1.041750	-3.944250
50%	0.000000	27.500000	0.000000	0.000000	1.000000	0.018500	-0.010000	-0.002000	3.741000	-1.281000
75%	0.000000	41.125000	0.625000	7.000000	2.250000	0.043000	0.014250	0.014250	5.875000	3.061000
max	0.210000	91.500000	12.000000	89.000000	20.000000	0.219000	0.106000	0.089000	25.595000	28.532000

```
In [20]: sns.pairplot(df)
```

```
Out[20]: <seaborn.axisgrid.PairGrid at 0x29b57b73588>
```

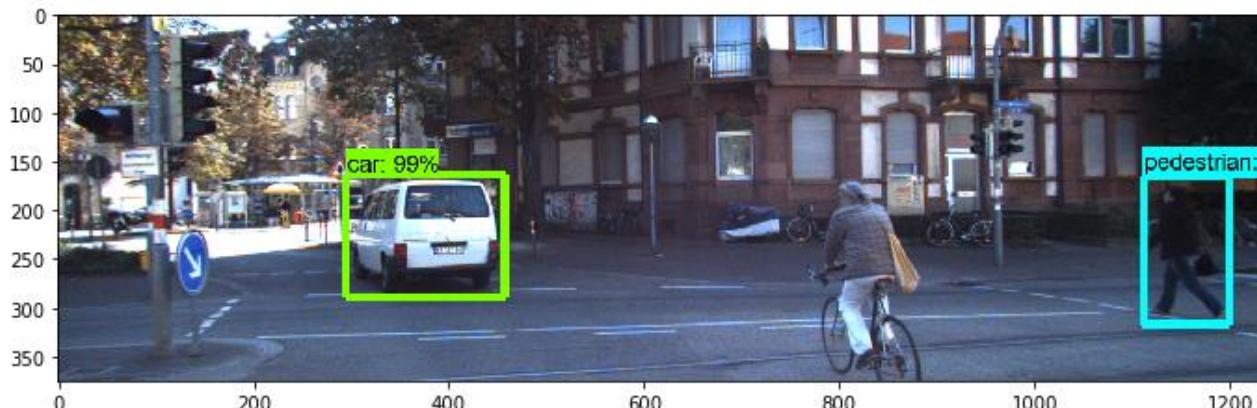


Pattern matching is done using Hungarian algorithm with the distance metrics:

$$d_{ij} = -\log(P(i \text{ and } j \text{ are same pattern})) = -\log(N(F_i - F_j | \mu_F, V_F))$$

If the probability that the patterns are same is near 1, the distance will be near zero. As the probability decreases, the distance increases. The log is required to compare small numbers without numerical issues.

Simple example



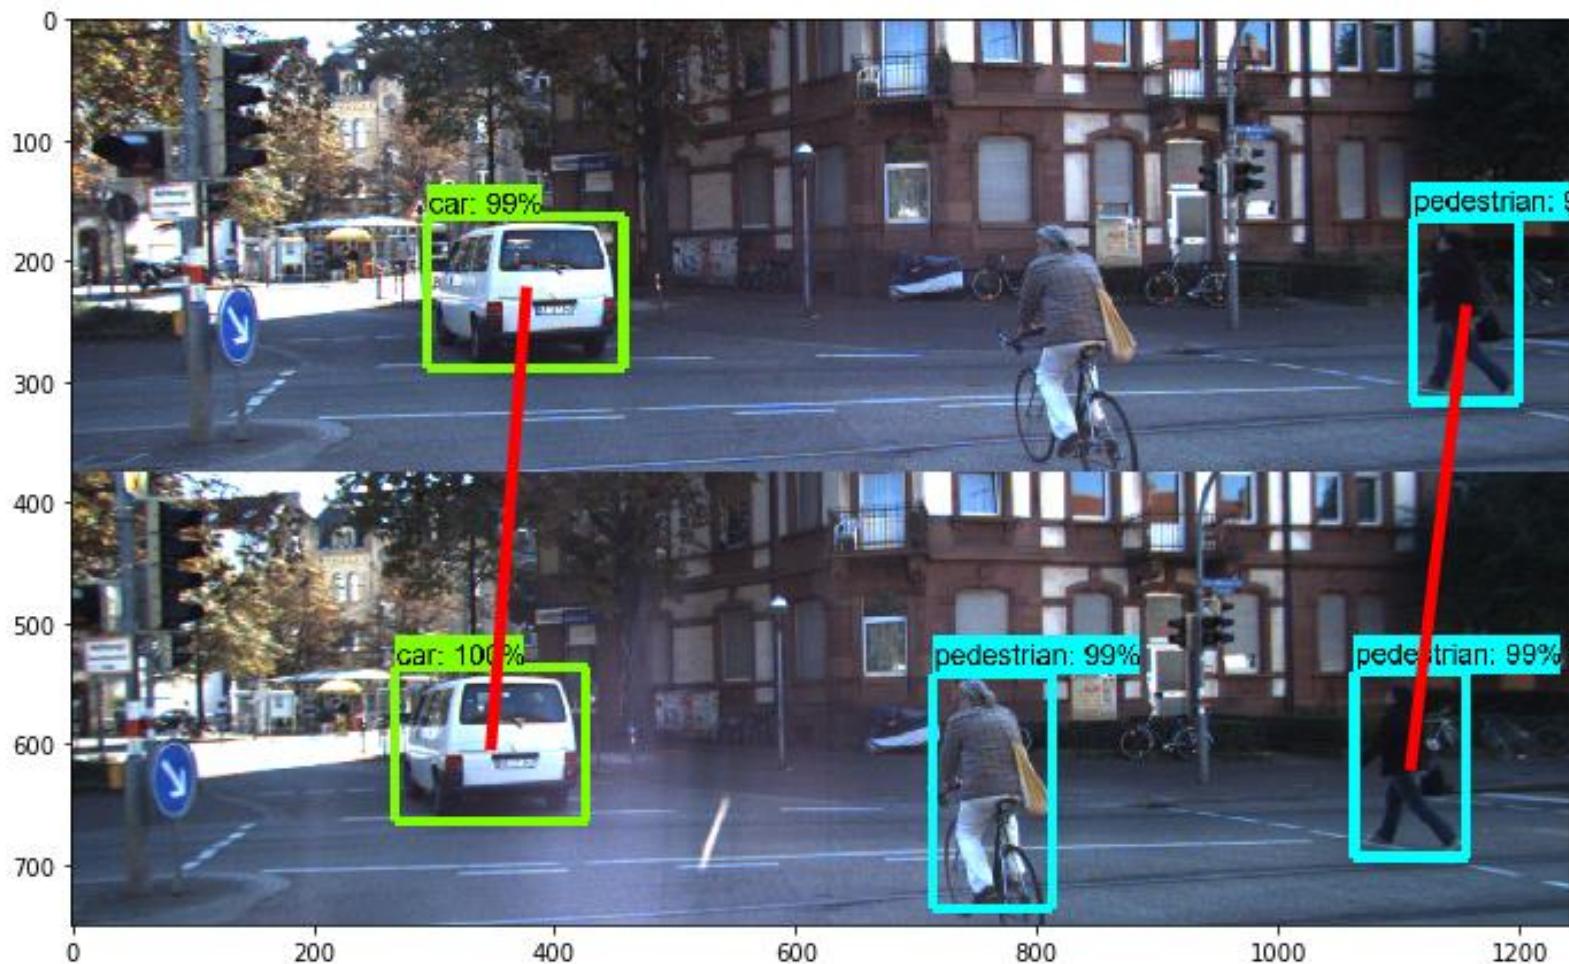
```
class,confidence,x,y,width,height,hue0,hue1,hue2,saturation,value
-----
2,1.00,1155.50,241.00,89.00,150.00,0.074,0.669,0.256,85.880,58.390
1,1.00,376.50,226.00,163.00,126.00,0.114,0.626,0.259,70.091,137.364
-----
1,1.00,347.50,225.50,159.00,127.00,0.104,0.614,0.282,66.123,148.577
2,1.00,1110.00,242.00,94.00,152.00,0.094,0.682,0.224,82.096,58.026
2,1.00,763.50,263.50,99.00,193.00,0.078,0.631,0.291,69.520,119.373
```

```
In [162]: np.set_printoptions(precision=0)
print(distance_matrix)
```

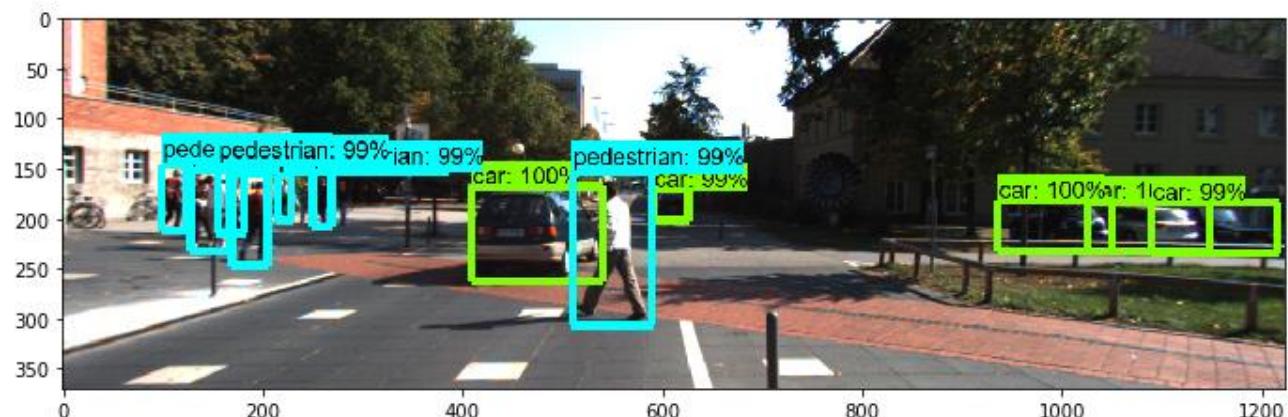
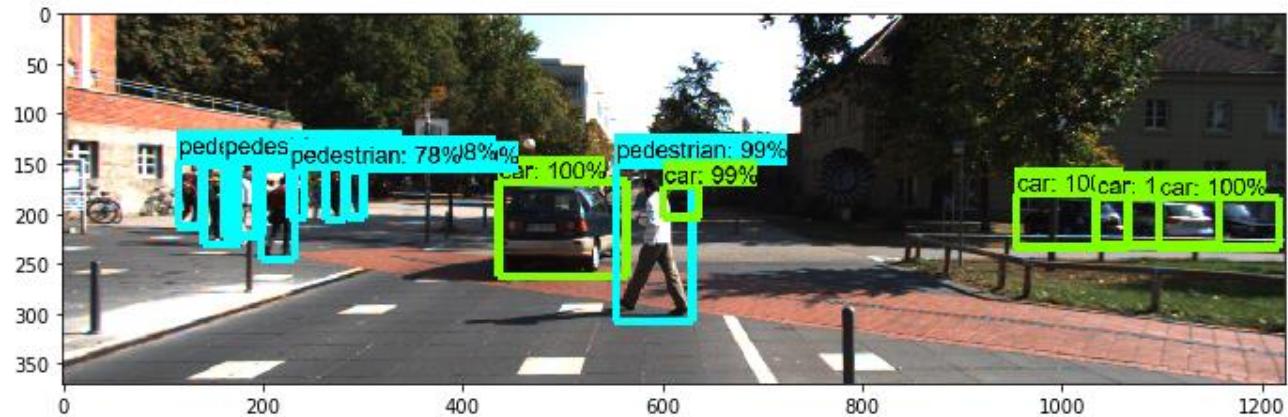
[[999. 6. 335.]	[6. 999. 571.]
------------------	-----------------

```
In [164]: row_ind, col_ind = linear_sum_assignment(distance_matrix)
print(row_ind)
print(col_ind)
```

[0 1]	[1 0]
-------	-------



Complicated example



```
class,confidence,x,y,width,height,hue0,hue1,hue2,saturation,value
-----
1,1.00,1010.50,206.50,113.00,53.00,0.354,0.438,0.208,74.793,48.389
1,1.00,1096.00,207.00,124.00,48.00,0.313,0.514,0.173,75.248,79.118
1,1.00,1157.50,208.50,121.00,49.00,0.274,0.551,0.175,79.360,81.186
1,1.00,500.50,215.00,129.00,94.00,0.367,0.389,0.244,75.836,70.045
2,1.00,592.00,229.50,78.00,155.00,0.550,0.254,0.196,59.824,104.288
1,1.00,619.50,187.50,37.00,31.00,0.286,0.323,0.391,58.958,58.623
2,1.00,159.00,190.50,36.00,79.00,0.434,0.220,0.346,77.767,116.533
2,1.00,293.50,180.50,17.00,45.00,0.354,0.299,0.346,78.097,70.915
2,1.00,127.50,180.00,25.00,64.00,0.493,0.174,0.333,97.595,144.548
2,1.00,216.00,200.00,34.00,90.00,0.422,0.256,0.322,82.717,112.057
2,1.00,184.50,184.00,33.00,72.00,0.413,0.229,0.358,68.688,159.507
2,1.00,177.50,185.50,33.00,75.00,0.395,0.244,0.361,70.105,148.906
2,0.98,272.50,179.50,19.00,51.00,0.397,0.261,0.342,80.279,69.196
2,0.79,235.00,180.50,16.00,47.00,0.483,0.209,0.309,94.992,127.202
-----
1,1.00,1087.00,207.50,124.00,49.00,0.292,0.507,0.201,71.784,76.482
1,1.00,993.50,207.00,115.00,50.00,0.348,0.410,0.242,72.279,47.679
1,1.00,475.00,215.00,132.00,94.00,0.311,0.366,0.323,69.133,74.729
1,1.00,609.00,186.00,36.00,30.00,0.144,0.298,0.557,59.536,30.091
2,1.00,550.50,230.00,81.00,152.00,0.564,0.262,0.175,60.986,116.519
2,1.00,187.50,200.50,37.00,91.00,0.394,0.285,0.322,78.999,107.080
2,1.00,147.50,191.50,39.00,81.00,0.381,0.260,0.359,69.986,128.833
2,1.00,222.50,178.50,15.00,45.00,0.376,0.330,0.293,78.055,120.033
2,1.00,260.00,180.50,20.00,55.00,0.394,0.267,0.339,79.143,69.303
2,1.00,112.00,178.50,26.00,65.00,0.509,0.172,0.319,90.011,143.473
1,1.00,1153.50,209.00,127.00,50.00,0.209,0.593,0.198,77.370,78.260
2,0.99,169.00,180.50,28.00,67.00,0.432,0.317,0.251,75.630,159.326
```

```

class,confidence,x,y,width,height,hue0,hue1,hue2,saturation,value
-----
1,1.00,1010.50,206.50,113.00,53.00,0.354,0.438,0.208,74.793,48.389
1,1.00,1096.00,207.00,124.00,48.00,0.313,0.514,0.173,75.248,79.118
1,1.00,1157.50,208.50,121.00,49.00,0.274,0.551,0.175,79.360,81.186
1,1.00,500.50,215.00,129.00,94.00,0.367,0.389,0.244,75.836,70.045
2,1.00,592.00,229.50,78.00,155.00,0.550,0.254,0.196,59.824,104.288
1,1.00,619.50,187.50,37.00,31.00,0.286,0.323,0.391,58.958,58.623
2,1.00,159.00,190.50,36.00,79.00,0.434,0.220,0.346,77.767,116.533
2,1.00,293.50,180.50,17.00,45.00,0.354,0.299,0.346,78.097,70.915
2,1.00,127.50,180.00,25.00,64.00,0.493,0.174,0.333,97.595,144.548
2,1.00,216.00,200.00,34.00,90.00,0.422,0.256,0.322,82.717,112.057
2,1.00,184.50,184.00,33.00,72.00,0.413,0.229,0.358,68.688,159.507
2,1.00,177.50,185.50,33.00,75.00,0.395,0.244,0.361,70.105,148.906
2,0.98,272.50,179.50,19.00,51.00,0.397,0.261,0.342,80.279,69.196
2,0.79,235.00,180.50,16.00,47.00,0.483,0.209,0.309,94.992,127.202
-----
1,1.00,1087.00,207.50,124.00,49.00,0.292,0.507,0.201,71.784,76.482
1,1.00,993.50,207.00,115.00,50.00,0.348,0.410,0.242,72.279,47.679
1,1.00,475.00,215.00,132.00,94.00,0.311,0.366,0.323,69.133,74.729
1,1.00,609.00,186.00,36.00,30.00,0.144,0.298,0.557,59.536,30.091
2,1.00,550.50,230.00,81.00,152.00,0.564,0.262,0.175,60.986,116.519
2,1.00,187.50,200.50,37.00,91.00,0.394,0.285,0.322,78.999,107.080
2,1.00,147.50,191.50,39.00,81.00,0.381,0.260,0.359,69.986,128.833
2,1.00,222.50,178.50,15.00,45.00,0.376,0.330,0.293,78.055,120.033
2,1.00,260.00,180.50,20.00,55.00,0.394,0.267,0.339,79.143,69.303
2,1.00,112.00,178.50,26.00,65.00,0.509,0.172,0.319,90.011,143.473
1,1.00,1153.50,209.00,127.00,50.00,0.209,0.593,0.198,77.370,78.260
2,0.99,169.00,180.50,28.00,67.00,0.432,0.317,0.251,75.630,159.326

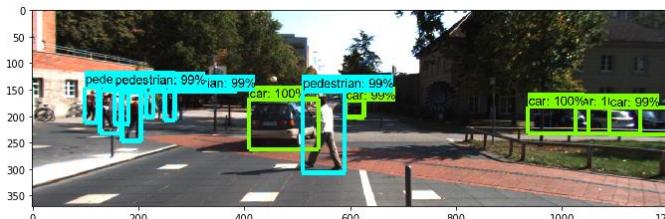
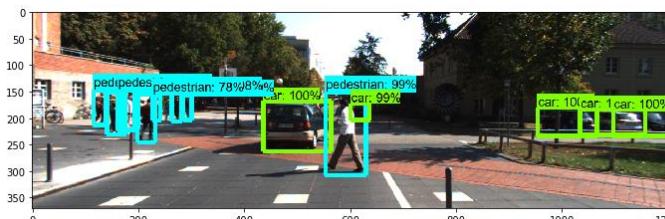
```

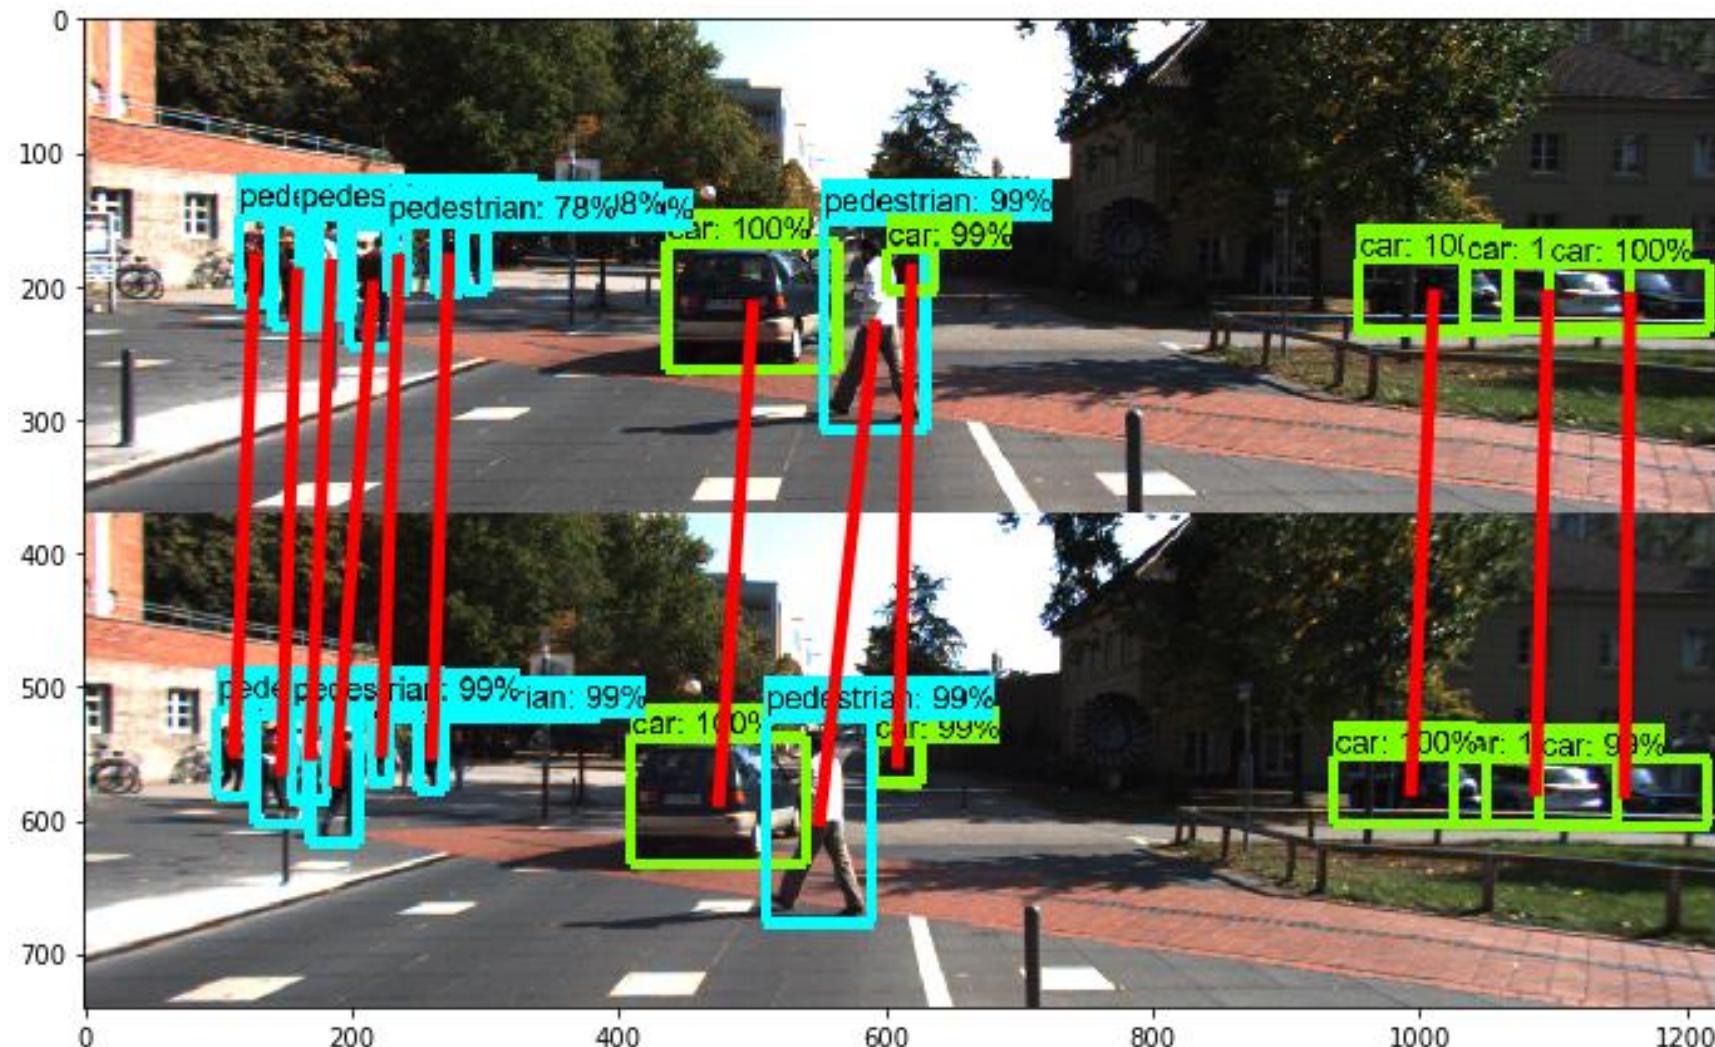
```
In [120]: np.set_printoptions(precision=0)
print(distance_matrix)
```

```
[[ 27.   6.  424.  315.  583.  999.  999.  999.  999.  999.  64.  999.]
 [ 6.   24.  543.  427.  695.  999.  999.  999.  999.  999.  21.  999.]
 [ 9.   46.  653.  532.  999.  999.  999.  999.  999.  999.  7.  999.]
 [ 566.  438.   7.  235.  154.  153.  257.  295.  232.  387.  693.  328.]
 [ 718.  616.  164.  600.   7.  372.  532.  690.  592.  999.  999.  651.]
 [ 453.  331.  191.   19.  542.  313.  345.  208.  155.  367.  588.  328.]
 [ 999.  999.  260.  406.  546.   29.   8.  54.   74.   35.  999.  31.]
 [ 999.  999.  246.  182.  631.  107.   93.  23.   9.  79.  999.  82.]
 [ 999.  999.  385.  449.  740.   96.   45.  40.   85.   6.  999.  25.]
 [ 999.  999.  157.  372.  372.   5.   22.  92.  92.   86.  999.  68.]
 [ 999.  999.  336.  443.  643.   79.   21.  53.  114.   35.  999.  14.]
 [ 999.  999.  314.  435.  613.   62.   14.  53.  106.   36.  999.  16.]
 [ 999.  999.  260.  206.  633.  102.   88.  25.   6.  68.  999.  79.]
 [ 999.  999.  327.  302.  717.  122.   86.  23.   49.   35.  999.  51.]]
```

```
In [122]: row_ind, col_ind = linear_sum_assignment(distance_matrix)
print(row_ind)
print(col_ind)
```

```
[ 0  1  2  3  4  5  6  8  9  10 12 13]
[ 1  0 10  2  4  3  6  9  5 11  8  7]
```



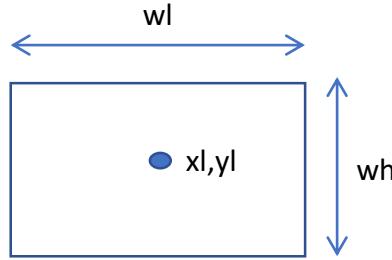


Too complicated!!! Curse of dimensionality makes probabilities (even with log) too small.


Simpler solution:

```
if (patterns_left[row].class_id == detections_right[col].class_id):
    c = abs(wl-wr)+abs(hl-hr)+abs(yl-yr) + STEREO_MATCHING_HORIZONTAL*abs(xl-xr)
else:
    c = 999999.0
if xl - xr < 0:
    c = 999999.0
cost[row][col] = c
```

Left camera pattern



Right camera detection



STEREO_MATCHING_HORIZONTAL = 0.4
Hungarian algorithm for matching.

Note: Patterns are not generated for the right camera. Patterns life spans in two cameras lead to complications.



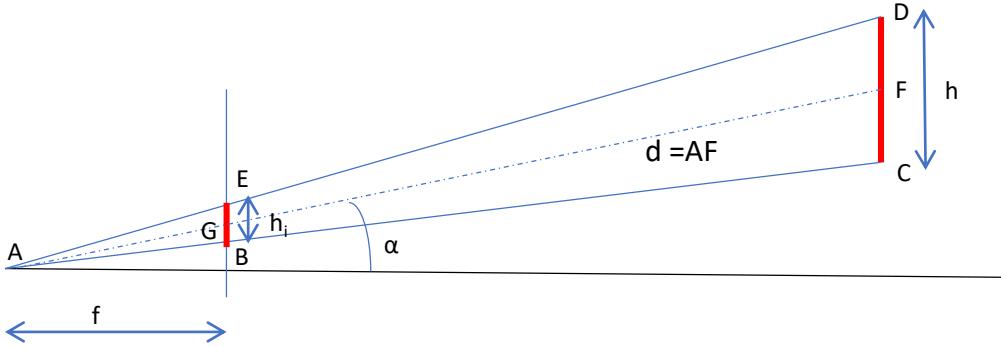
Not perfect, but good enough for the moment.

Research topic for 2019



After implementing stereo vision, we have two distance estimates:

1. Distance based on stereo vision. Accurate in short distances (for Kitti, 20 meters)
2. Distance based on size. Can be used in long distances where stereo vision is inaccurate.



$$d_{size} = \frac{f * r}{\cos(\alpha) * \cos(\beta) * r_i * s_h / p_h}$$

s_h = sensor height (m)
 p_h = image height (pixels)
 r_i = pattern radius (pixels)
 r = body radius (m), mean from class specific distribution
 f = focal length (m)
 α = altitude (rad)
 β = azimuth (rad)

$$d_{stereo} = \frac{f * b}{\cos(\alpha) * \cos(\beta) * ds * s_w / p_w}$$

s_w = sensor width (m)
 p_w = image width (pixels)
 f = focal length (m)
 b = base line (m)
 ds = disparity (pixels)
 α = altitude (rad)
 β = azimuth (rad)

Combining distance estimates:

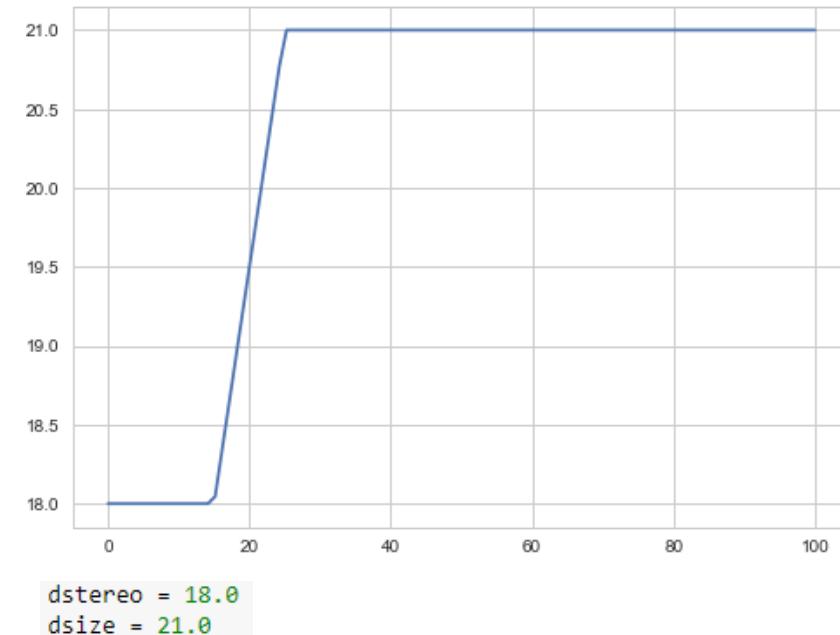
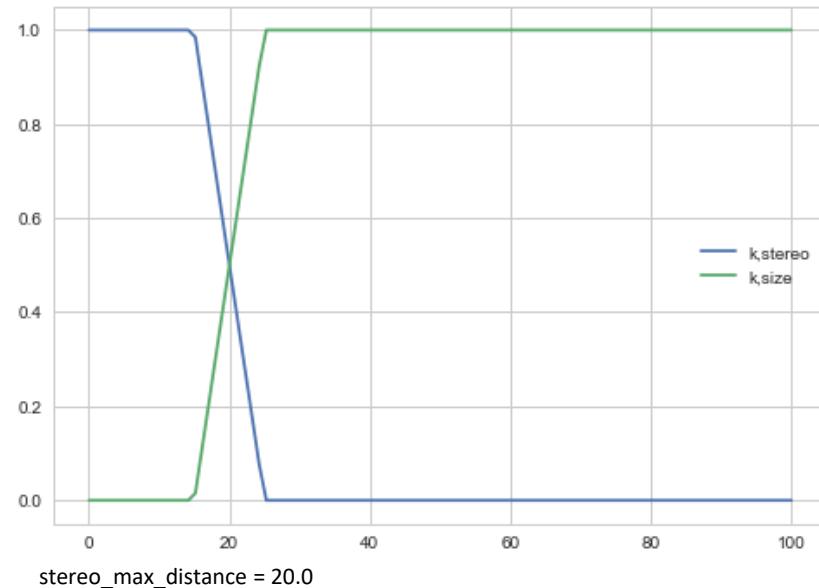
$$d = k_{\text{stereo}} * d_{\text{stereo}} + k_{\text{size}} * d_{\text{size}}$$

```
fraction = 0.25
def calculate_coefficients(estimated_distance, stereo_max_distance):
    if estimated_distance < (1-fraction)*stereo_max_distance:
        k_size = 0.0
        k_stereo = 1.0
    elif estimated_distance > (1+fraction)*stereo_max_distance:
        k_size = 1.0
        k_stereo = 0.0
    else:
        l1 = estimated_distance - (1-fraction)*stereo_max_distance
        l2 = (1+fraction)*stereo_max_distance - (1-fraction)*stereo_max_distance
        k_size = 11/l2
        k_stereo = 1 - k_size
    return k_stereo, k_size
```

Initialization:

$$\text{estimated_distance} = 0.5 * d_{\text{stereo}} + 0.5 * d_{\text{size}}$$

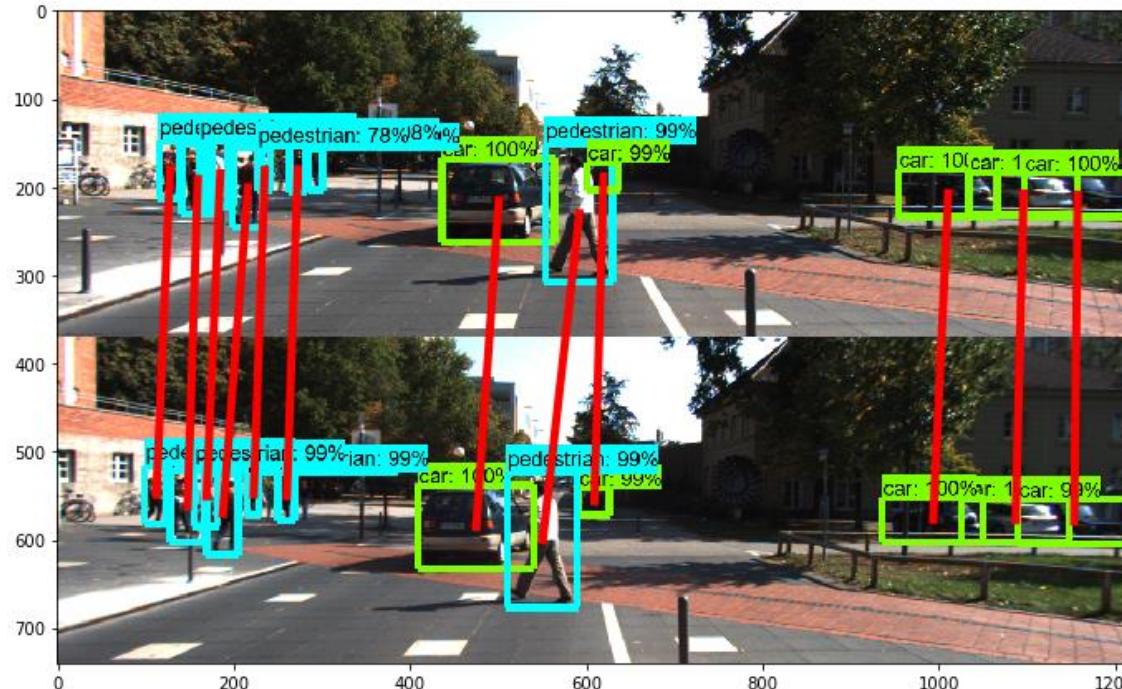
Procedure is iterated until convergence or max_iter
(or just used once?)



Estimating disparity using matched patterns


Left camera pattern:
x_min_left
x_max_left

Right camera pattern:
x_min_right
x_max_right



$$\text{Pattern disparity} = 0.5 * (\text{x_min_left} + \text{x_max_left}) - 0.5 * (\text{x_min_right} + \text{x_max_right})$$

3D projection



$$(x_b, y_b, z_b) = t^* (x_p, y_p, z_p)$$

Where:

$$(x_p, y_p, z_p) = \left(-\frac{s_w}{2} + p_x * \frac{s_w}{p_w}, \frac{s_h}{2} - p_y * \frac{s_h}{p_h}, -f \right)$$

$$t = \frac{d}{\sqrt{x_p^2 + y_p^2 + z_p^2}}$$

s_w = sensor width (m)

s_h = sensor height (m)

p_w = image width (pixels)

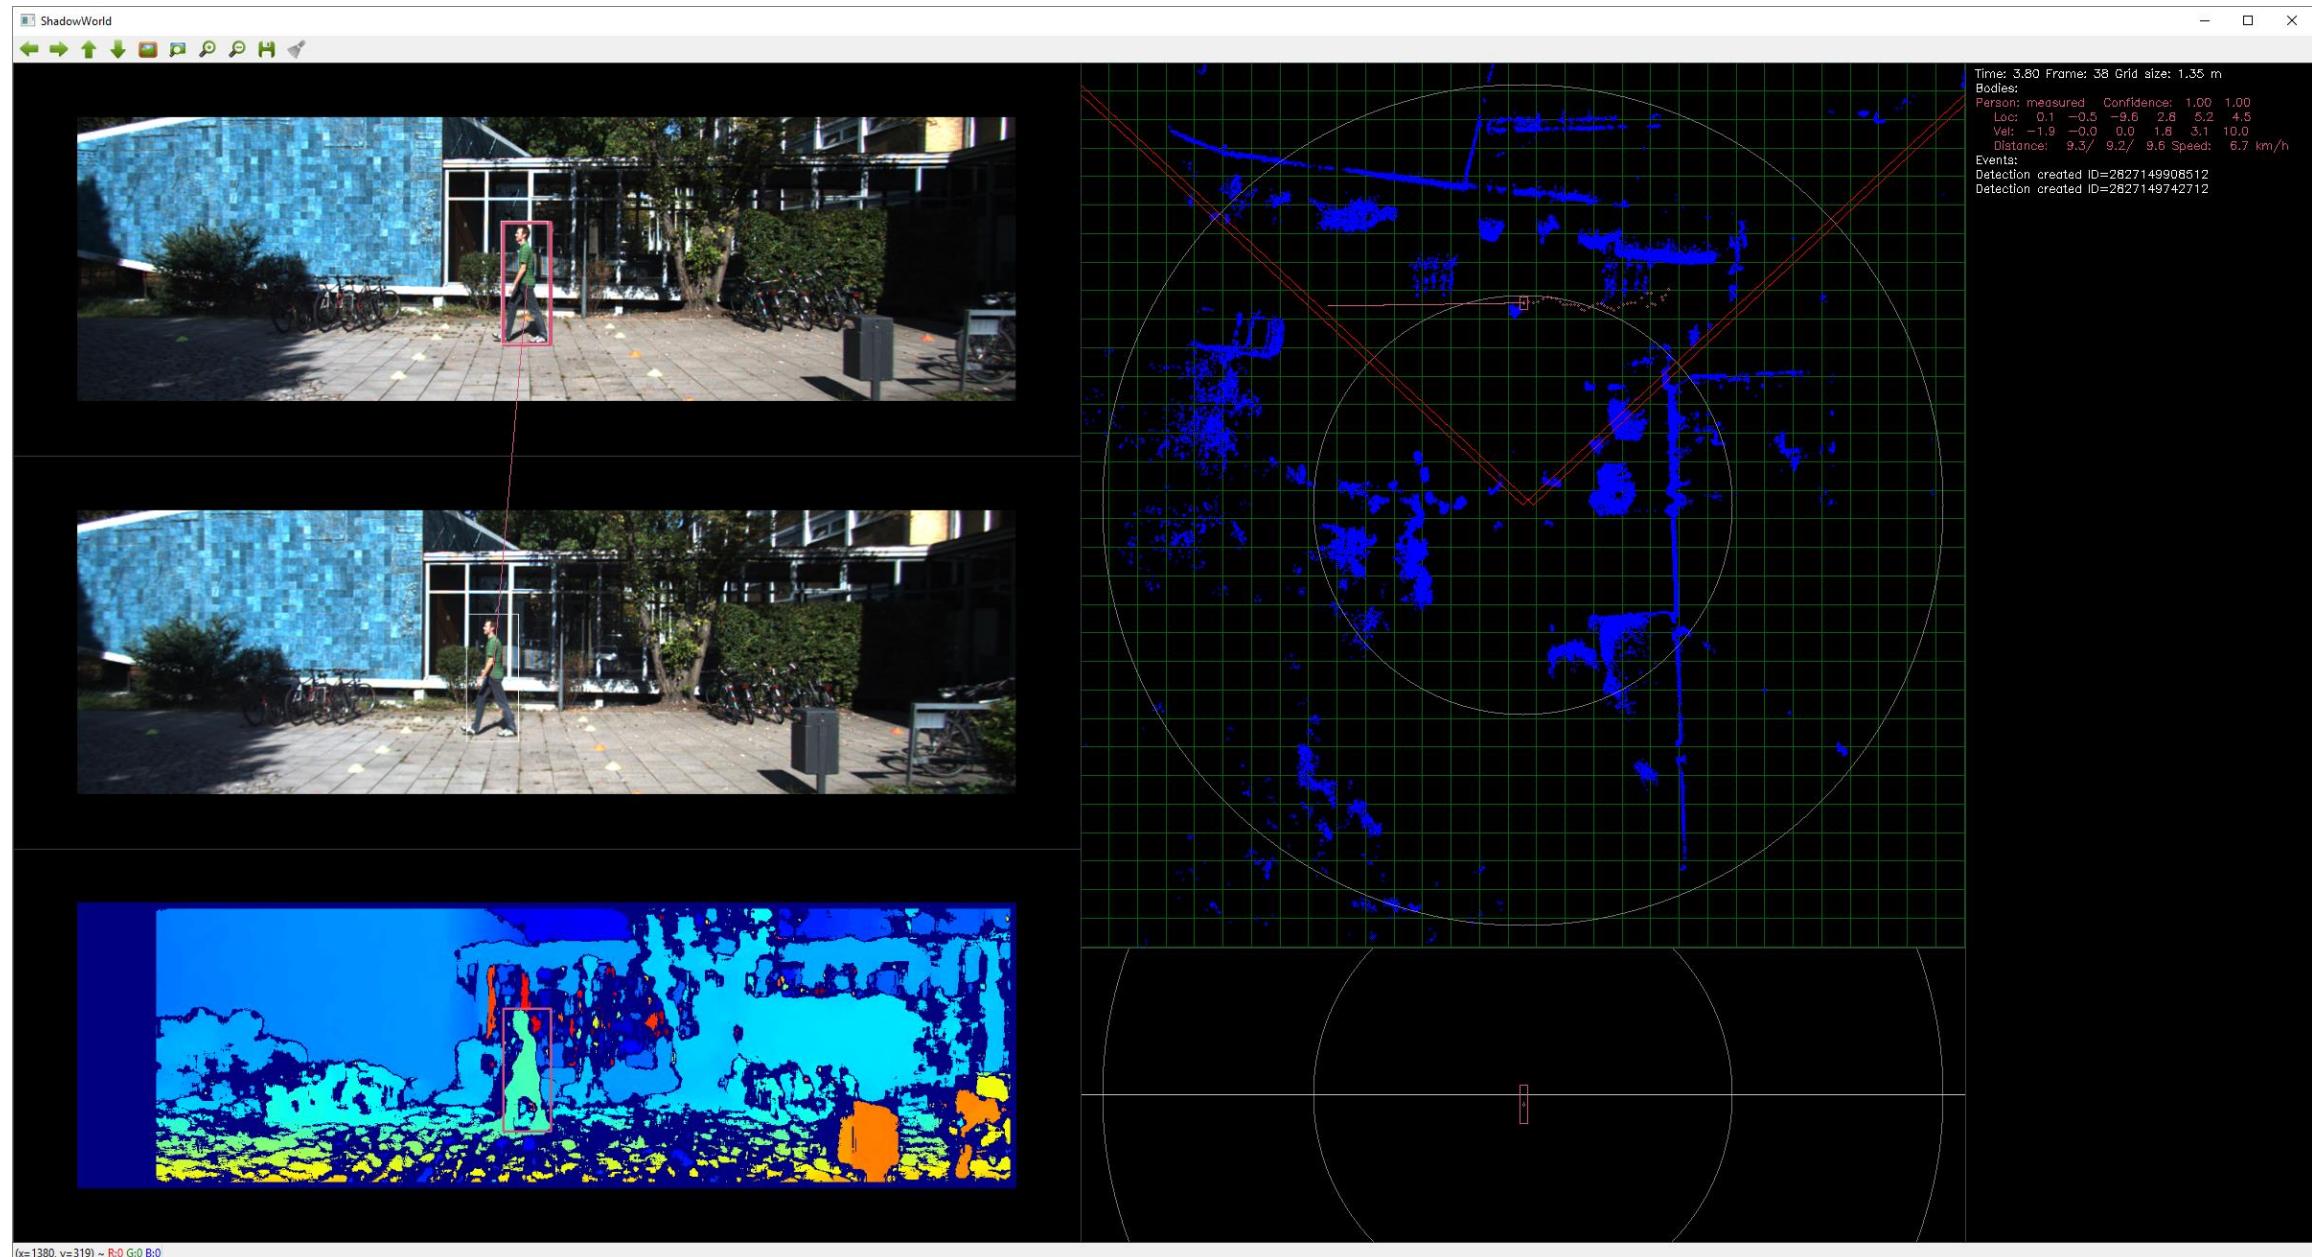
p_h = image height (pixels)

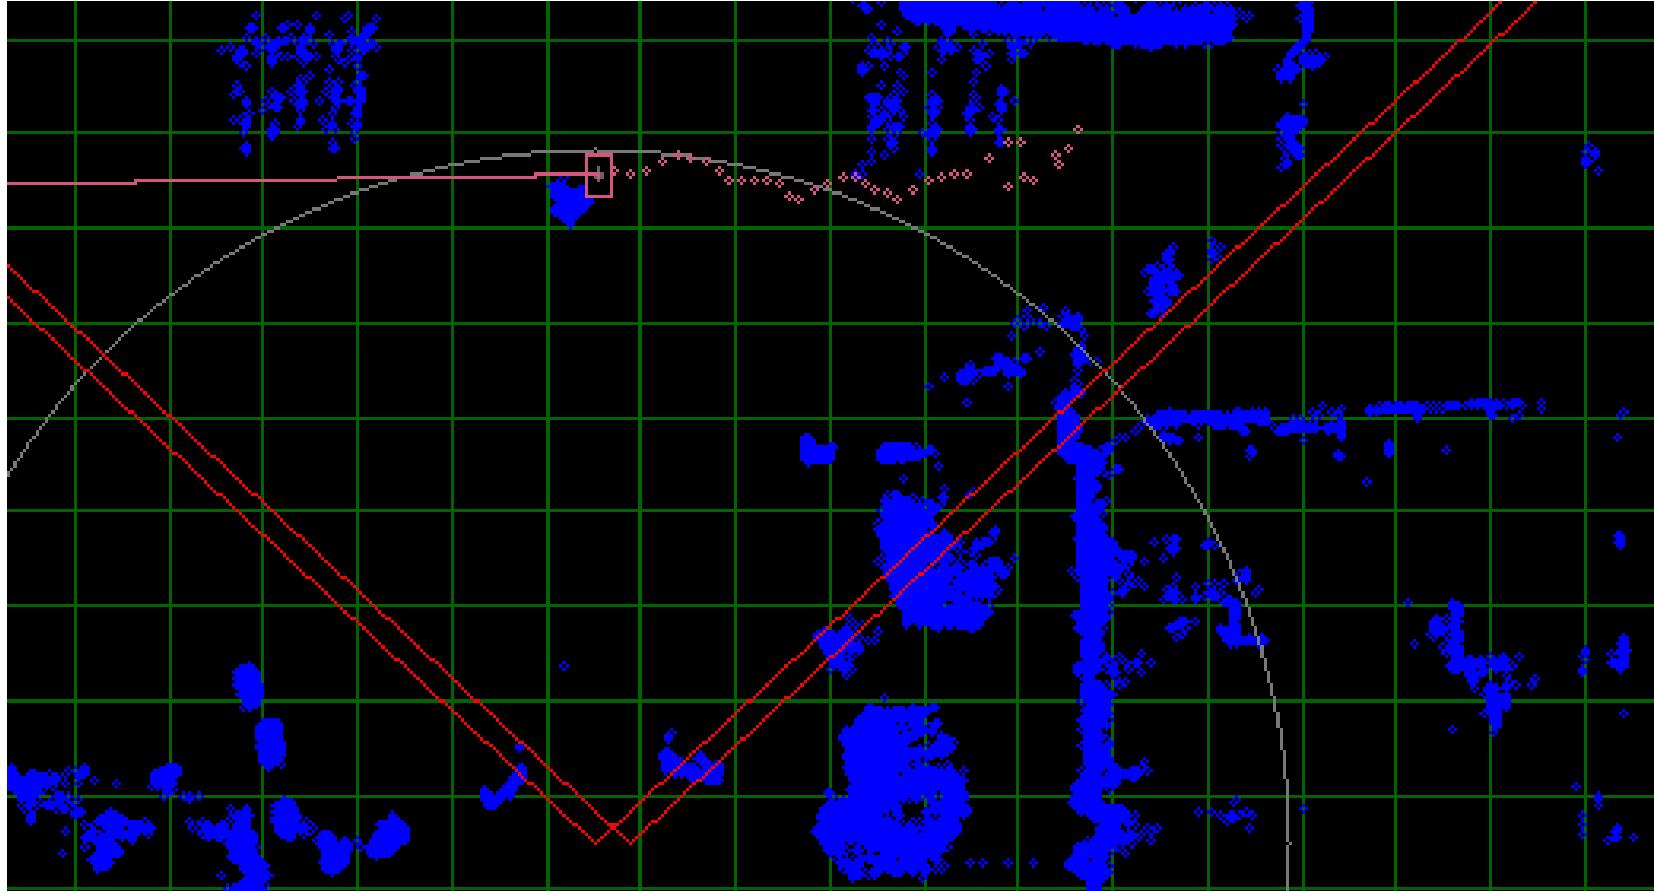
f = focal length (m)

p_x = pattern center point location (x, pixels)

p_y = pattern center point location (y, pixels)

Note! Only left image used. Right image is used only for disparity calculation (in the context of distance estimation and 3D projection).





Localizations follows Velodyne cloud pretty well!

Current overall matching algorithm:

- For left camera:
 - Create detections
 - Predict pattern locations
 - Match detections to patterns
 - For each unmatched detection
 - Create pattern
 - Create body
- For right camera:
 - Create detections
 - Match right camera detections to left camera patterns
 - Predict body locations
 - Create location measurements based on size and stereo
 - Correct body locations

Note: Pattern and body removals are not described.

Research topic!

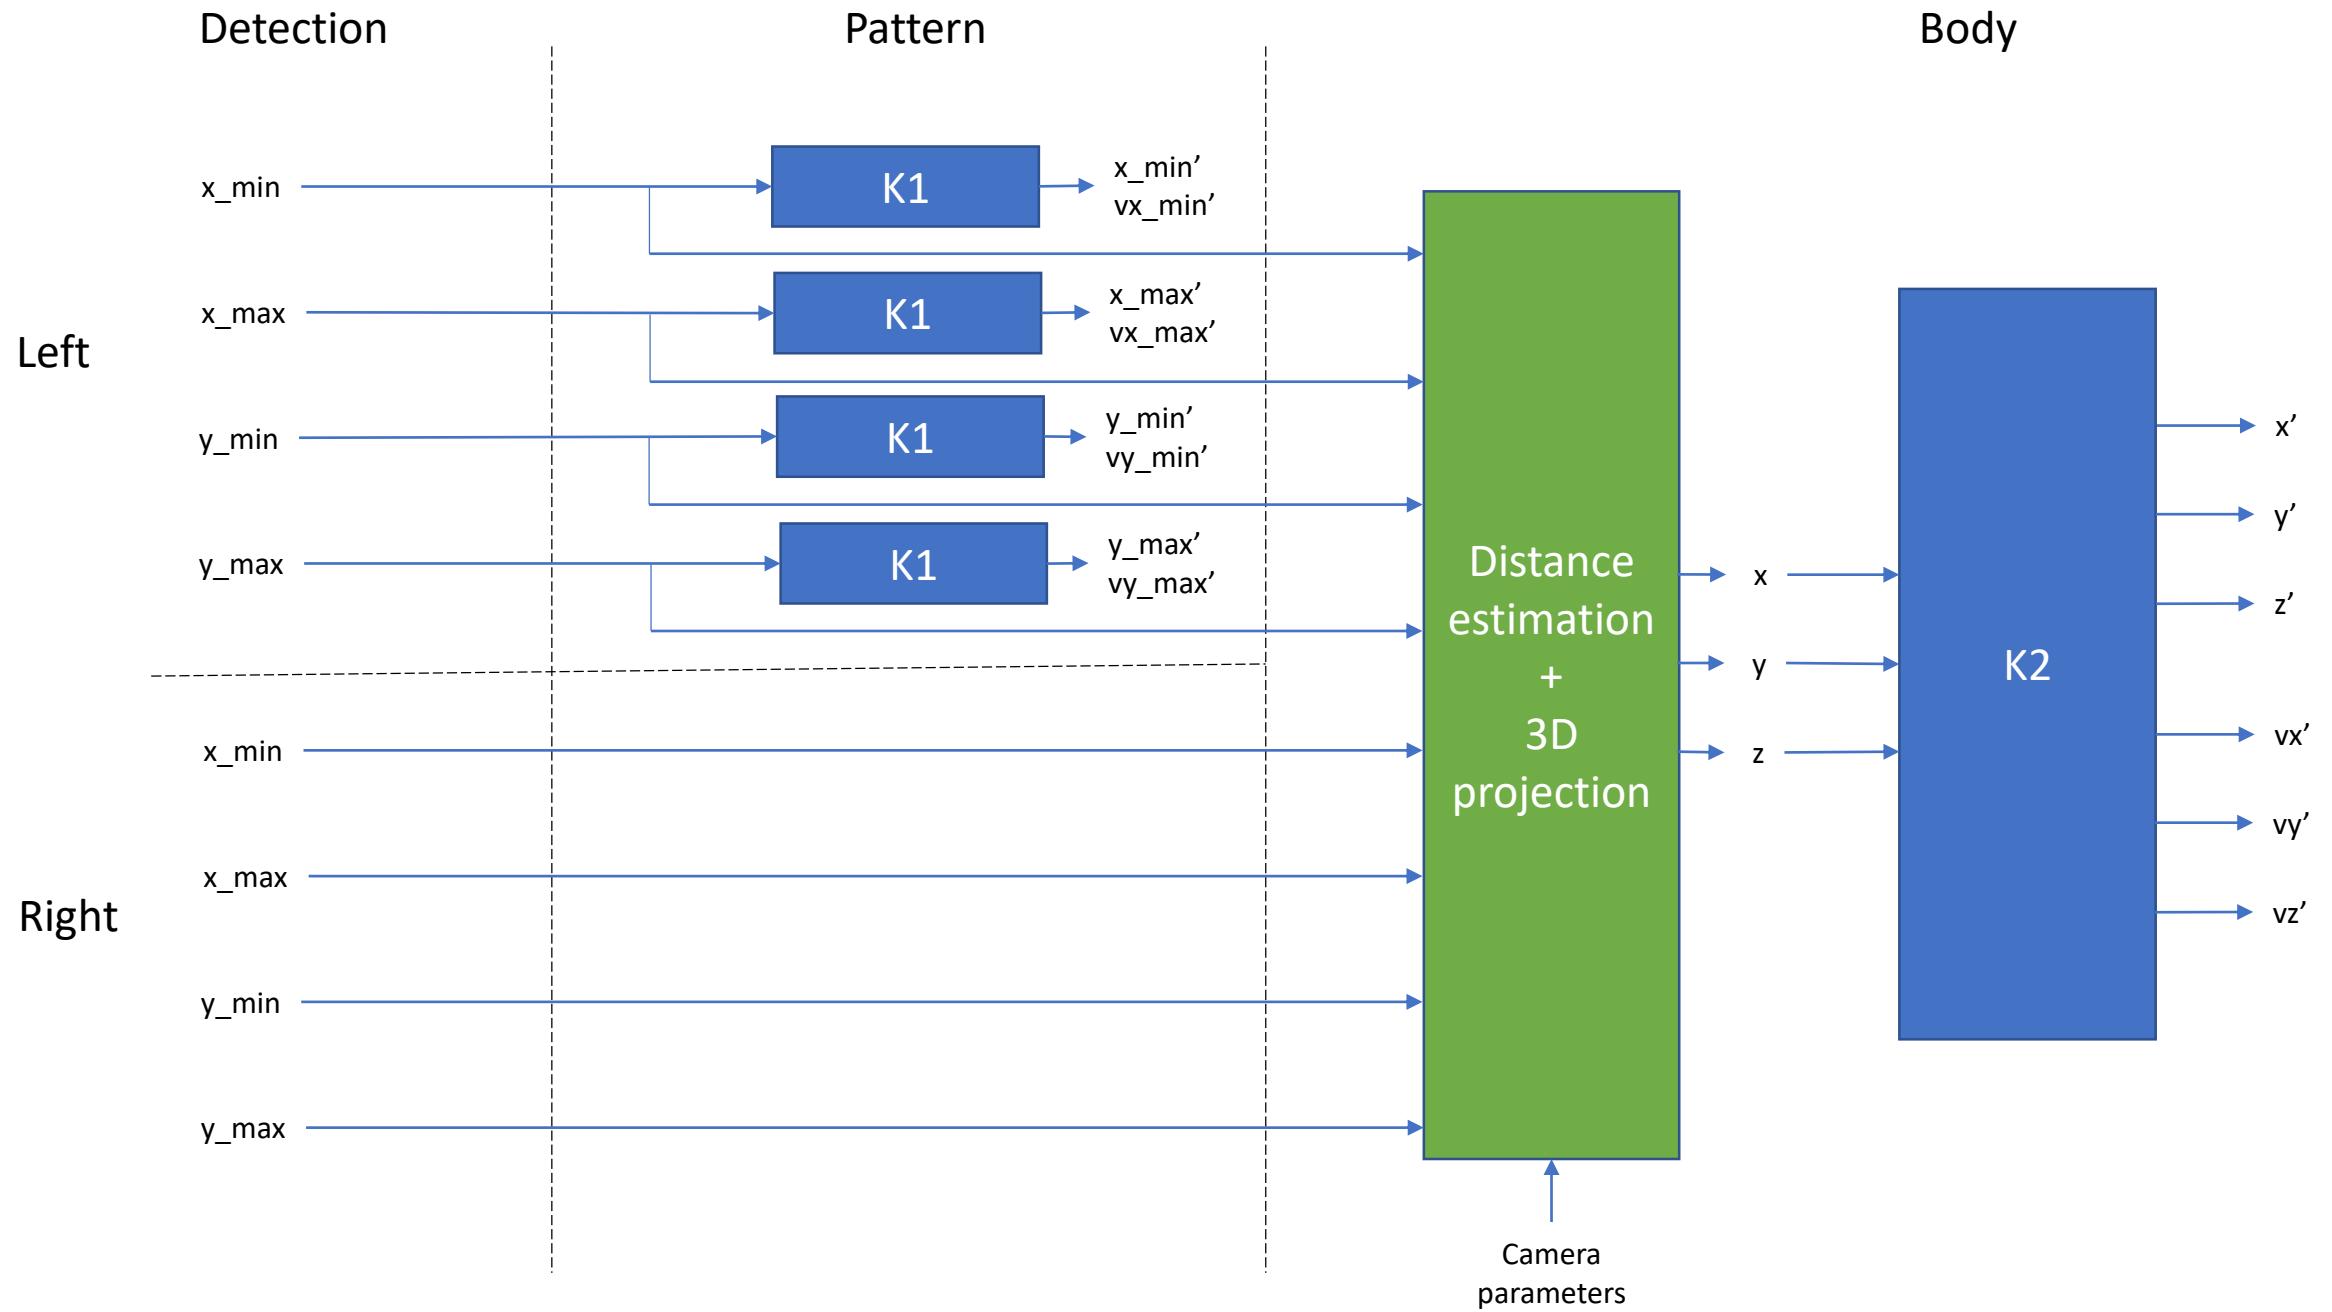
The joint estimation of pattern and body state would be a good area for further study. The current algorithm is by no means optimal.



Kalman Filter Parameter Adjustments

Two levels of filtering:

1. Pattern
 - Pattern movement estimation for missing detections
2. Body
 - Coordinate and velocity estimation



Before

```
PATTERN_ALFA = 200.0 # Pattern initial location error variance
PATTERN_BETA = 10000.0 # Pattern initial velocity error variance
PATTERN_C = np.array([[1.0, 0.0]]) # Pattern measurement matrix
PATTERN_Q = np.array([200.0]) # Pattern measurement variance
PATTERN_R = np.array([[0.1, 0.0],
                     [0.0, 1.0]]) # Pattern state equation covariance
"
```

```
BODY_ALFA = 100000.0 # Body initial location error variance
BODY_BETA = 100000.0 # Body initial velocity error variance
BODY_C = np.array([[1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 1.0, 0.0, 0.0, 0.0, 0.0]
                 ]) # Body measurement matrix
BODY_DATA_COLLECTION_COUNT = 30 # How many frames until notification
BODY_Q = np.array([[200.0, 0.0, 0.0],
                  [0.0, 200.0, 0.0],
                  [0.0, 0.0, 200.0]]) # Body measurement variance 200
BODY_R = np.array([[0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.0, 0.0, 0.0]
                 ]) # Body state equation covariance
"
```

After

```
PATTERN_ALFA = 100.0 # Pattern initial location error variance
PATTERN_BETA = 10000.0 # Pattern initial velocity error variance
PATTERN_C = np.array([[1.0, 0.0]]) # Pattern measurement matrix
PATTERN_Q = np.array([100.0]) # Pattern measurement variance
PATTERN_R = np.array([[0.0, 0.0],
                     [0.0, 100.0]]) # Pattern state equation covariance
"
```

```
BODY_ALFA = 10.0 # Body initial location error variance
BODY_BETA = 100.0 # Body initial velocity error variance
BODY_C = np.array([[1.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 1.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 1.0, 0.0, 0.0, 0.0]
                 ]) # Body measurement matrix
BODY_DATA_COLLECTION_COUNT = 30 # How many frames until notification
BODY_Q = np.array([[25.0, 0.0, 0.0],
                  [0.0, 25.0, 0.0],
                  [0.0, 0.0, 25.0]]) # Body measurement variance 200
BODY_R = np.array([[0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 1.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.1, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.0, 0.0, 1.0]
                 ]) # Body state equation covariance
"
```

```

PATTERN_ALFA = 100.0 # Pattern initial location error variance
PATTERN_BETA = 10000.0 # Pattern initial velocity error variance
PATTERN_C = np.array([[1.0, 0.0]]) # Pattern measurement matrix
PATTERN_Q = np.array([100.0]) # Pattern measurement variance
PATTERN_R = np.array([[0.0, 0.0],
                     [0.0, 100.0]]) # Pattern state equation covariance

```

```

BODY_ALFA = 10.0 # Body initial location error variance
BODY_BETA = 100.0 # Body initial velocity error variance
BODY_C = np.array([[1.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 1.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 1.0, 0.0, 0.0, 0.0]
                 ]) # Body measurement matrix
BODY_DATA_COLLECTION_COUNT = 30 # How many frames until notification
BODY_Q = np.array([[25.0, 0.0, 0.0],
                  [0.0, 25.0, 0.0],
                  [0.0, 0.0, 25.0]]) # Body measurement variance 200
BODY_R = np.array([[0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 1.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.1, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
                  [0.0, 0.0, 0.0, 0.0, 0.0, 1.0]
                 ]) # Body state equation covariance
"
```

PATTERN_ALFA

- $2\text{std} = 20\text{px}$, first detection bounding box accuracy

PATTERN_BETA

- $2\text{std} = 200 \text{ px/s} = 20 \text{ px / frame}$

PATTERN_Q

- $2\text{std} = 20\text{px}$, detection accuracy

PATTERN_R

- Location, no teleporting
- Velocity, $2\text{std} = 20\text{px/frame} = 200 \text{ px/s}$

BODY_ALFA

- $2\text{std} = 6 \text{ m}$, first location accuracy

BODY_BETA

- $2\text{std} = 20 \text{ m/s}$

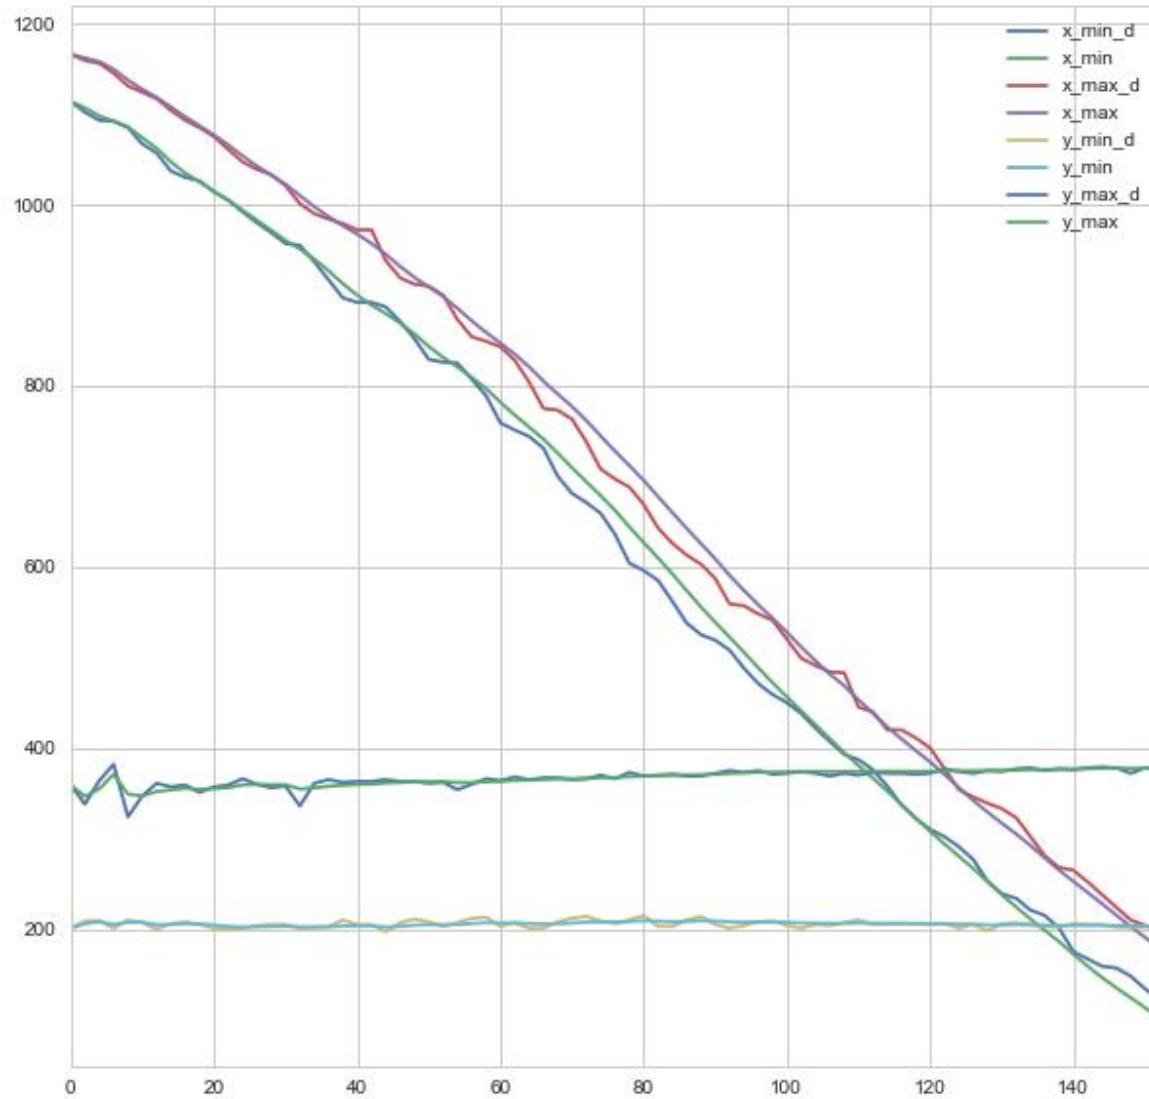
BODY_Q

- $2\text{std} = 10\text{m}$, measurement accuracy

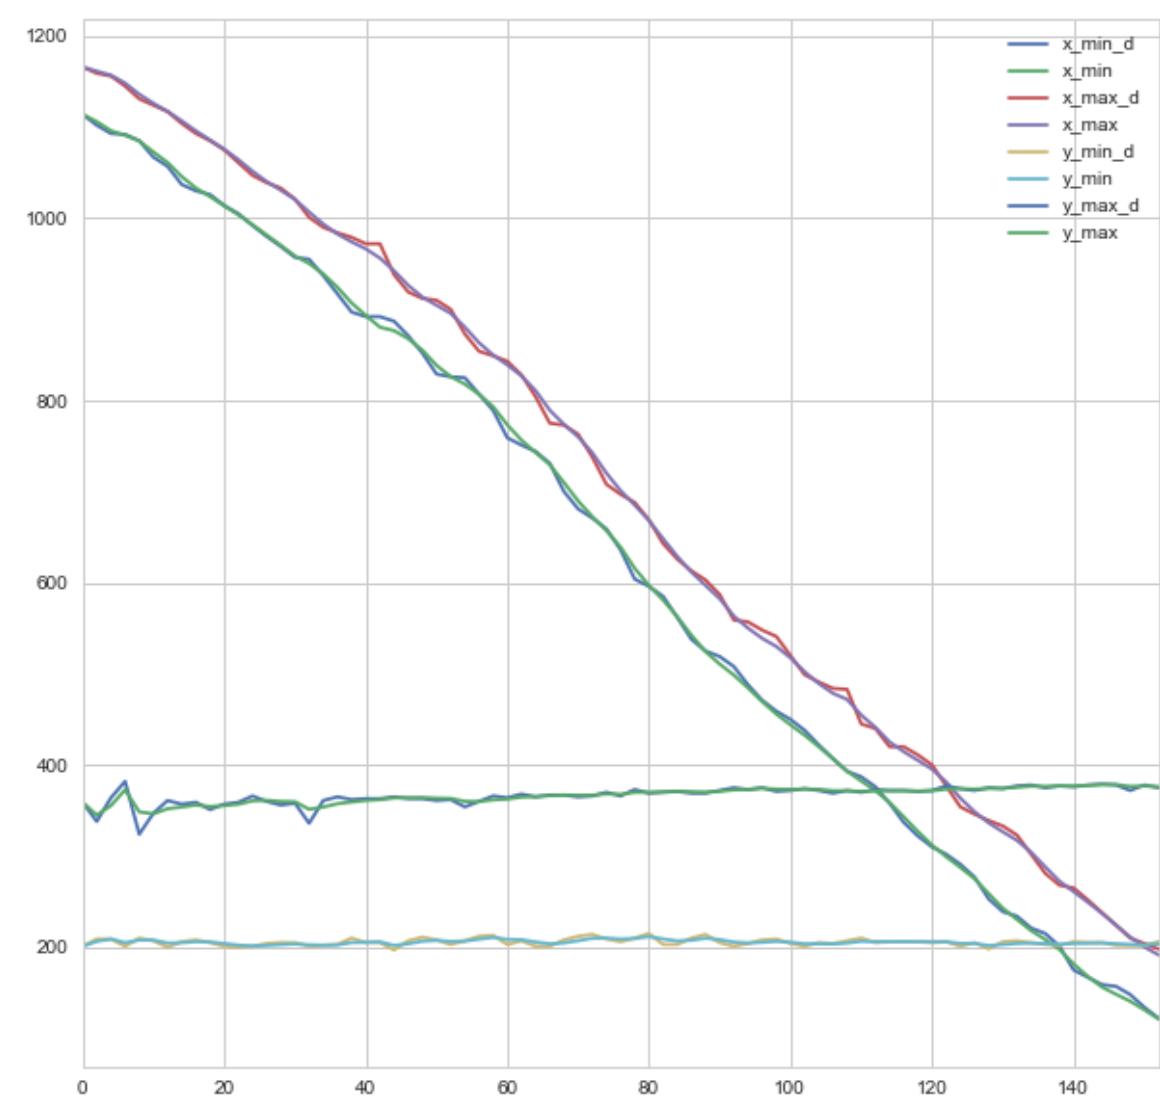
BODY_R

- Location, no teleporting
- Velocity (x,z), $2\text{std} = 2 \text{ m/s}^2$
- Velocity (y), $2\text{std} = 0,6 \text{ m/s}^2$

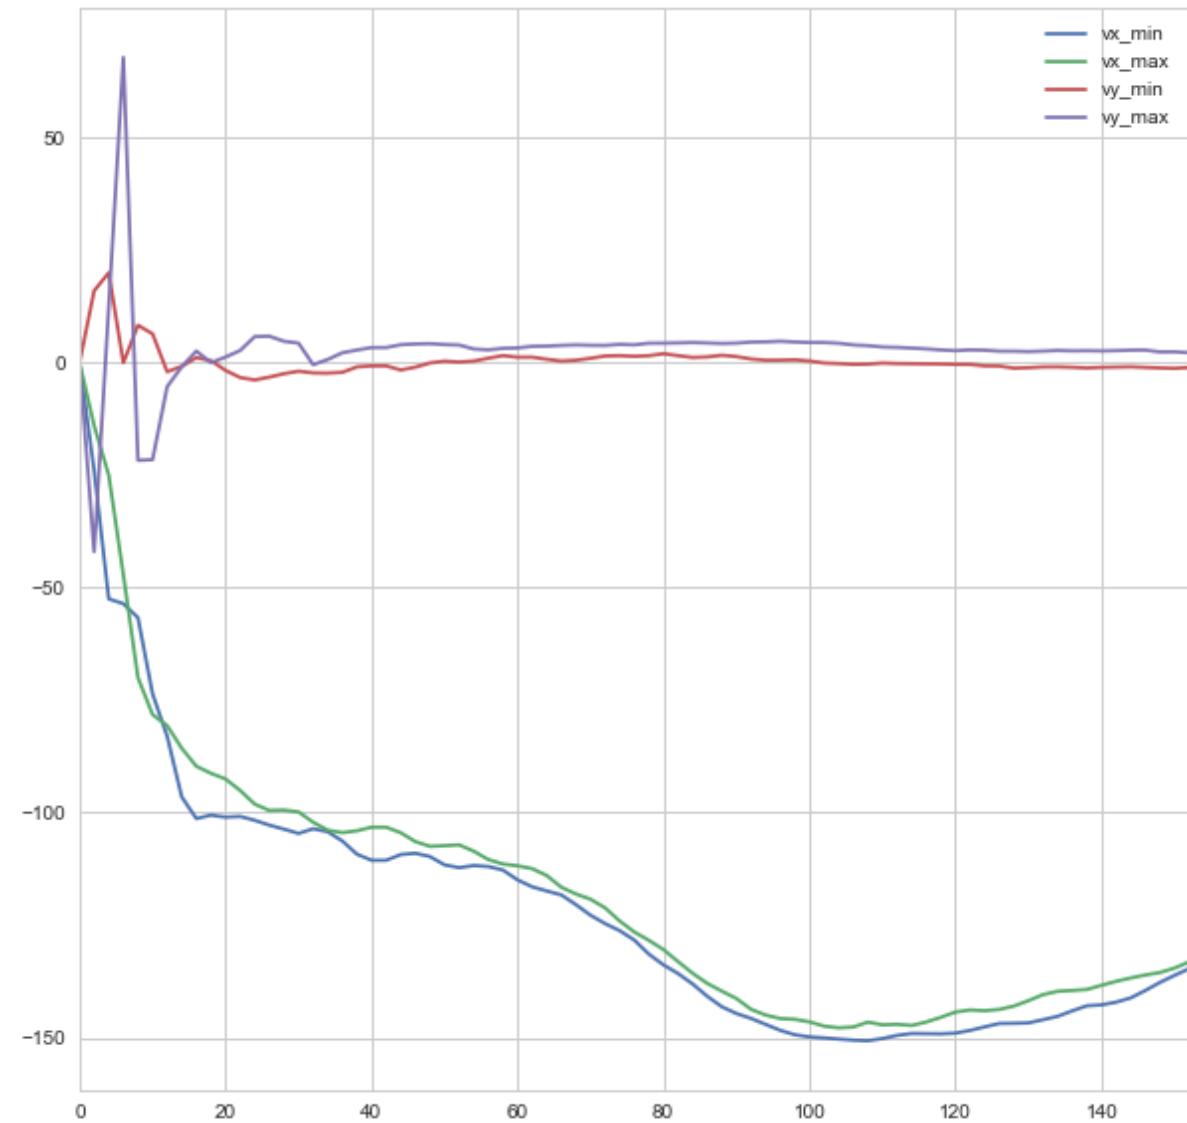
Pattern location / before



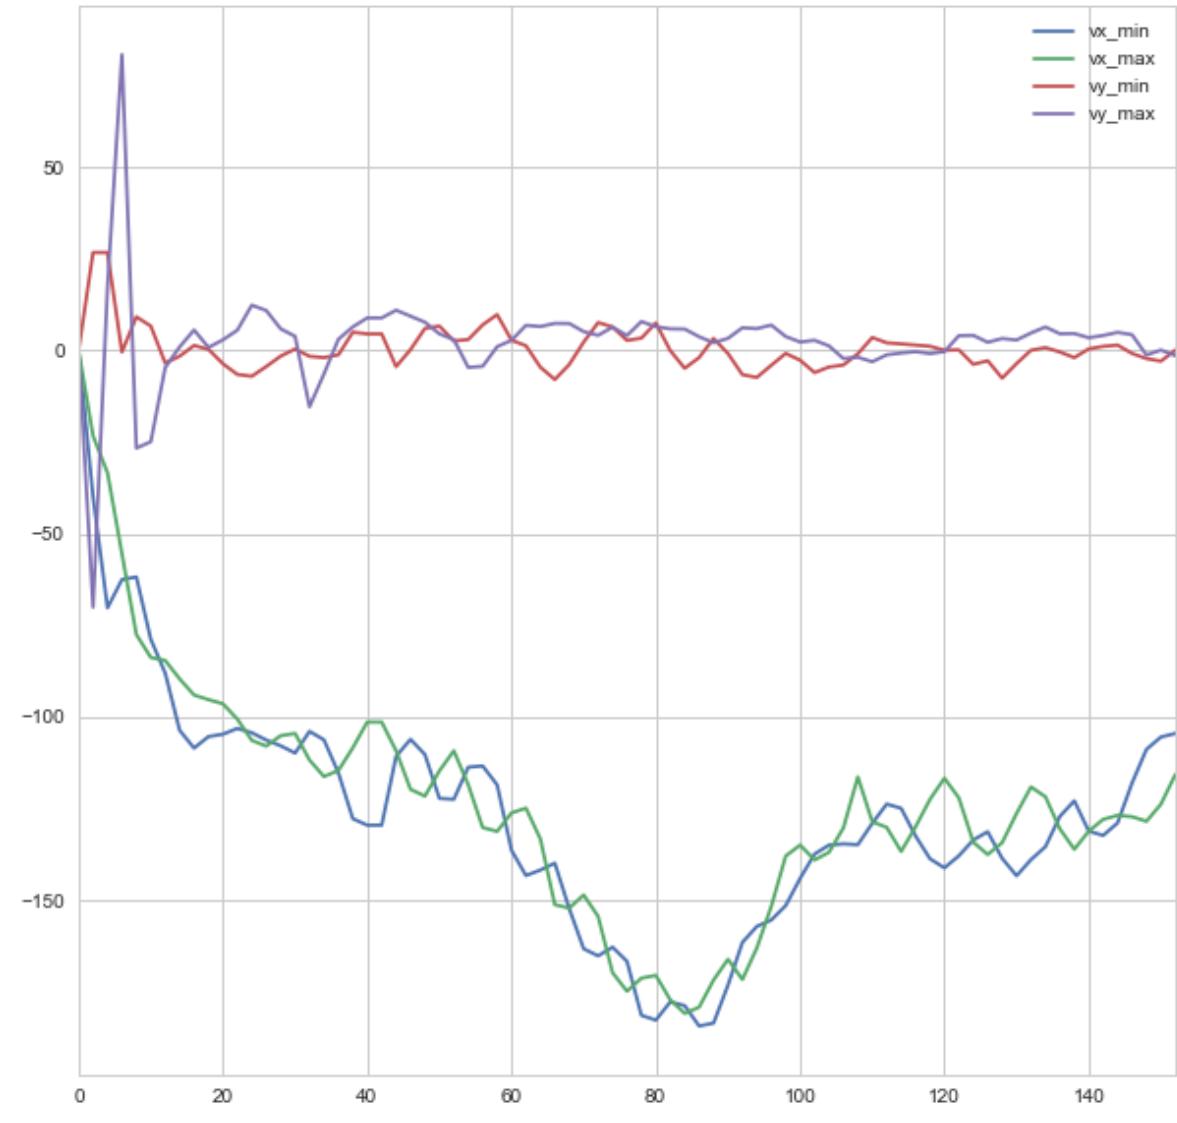
Pattern location / after



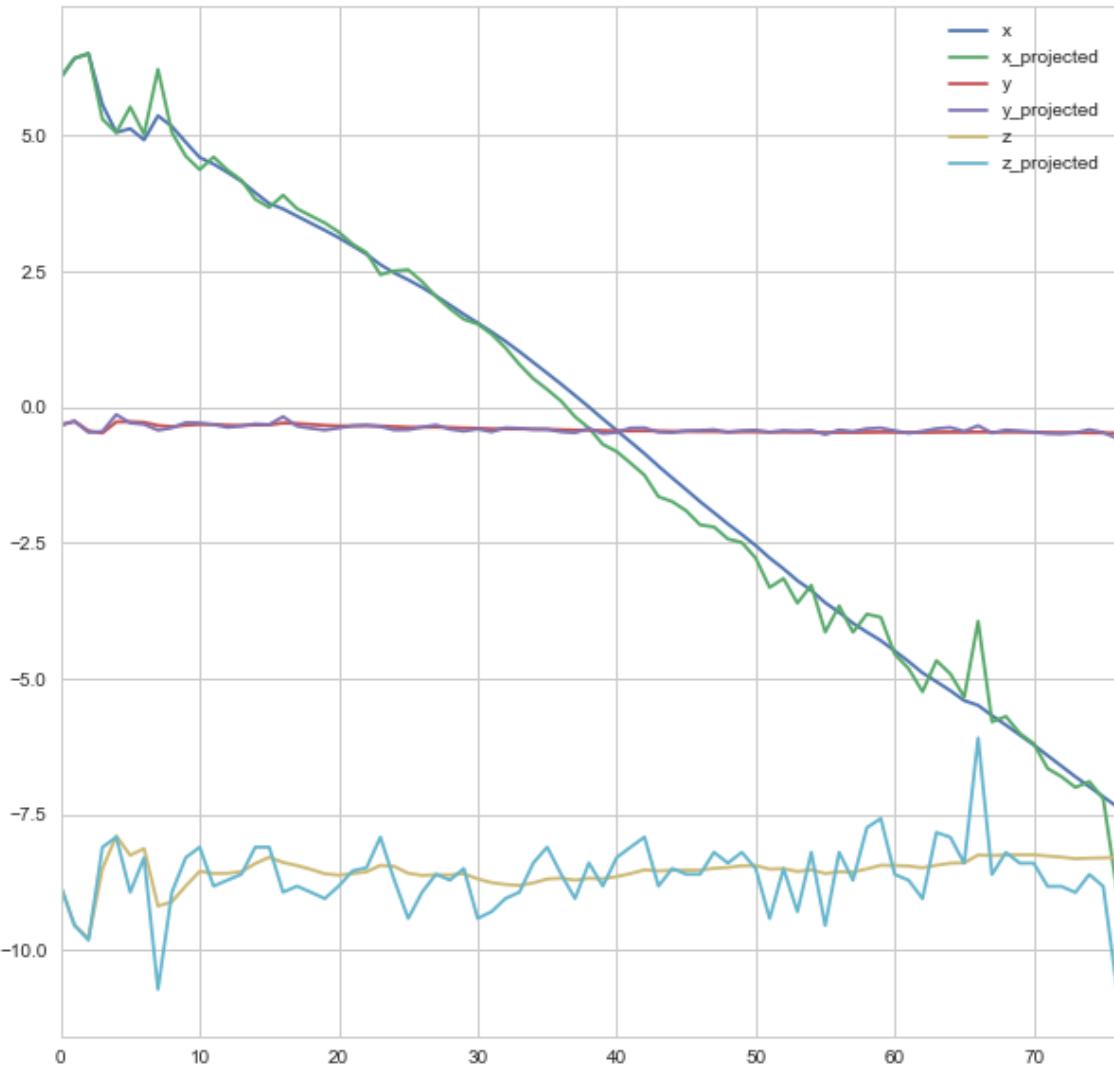
Pattern velocity / before




Pattern location / after



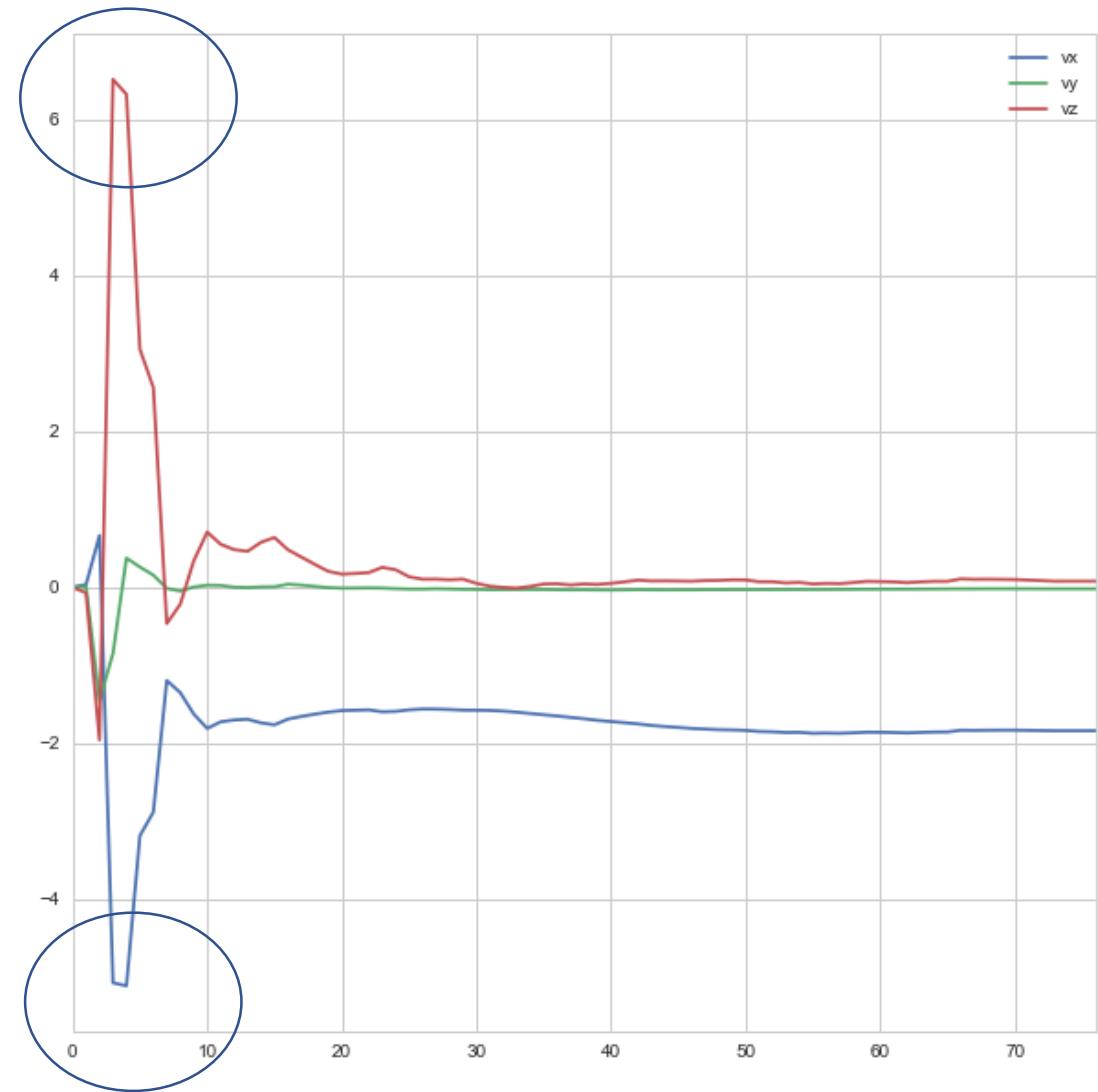
Body location / before



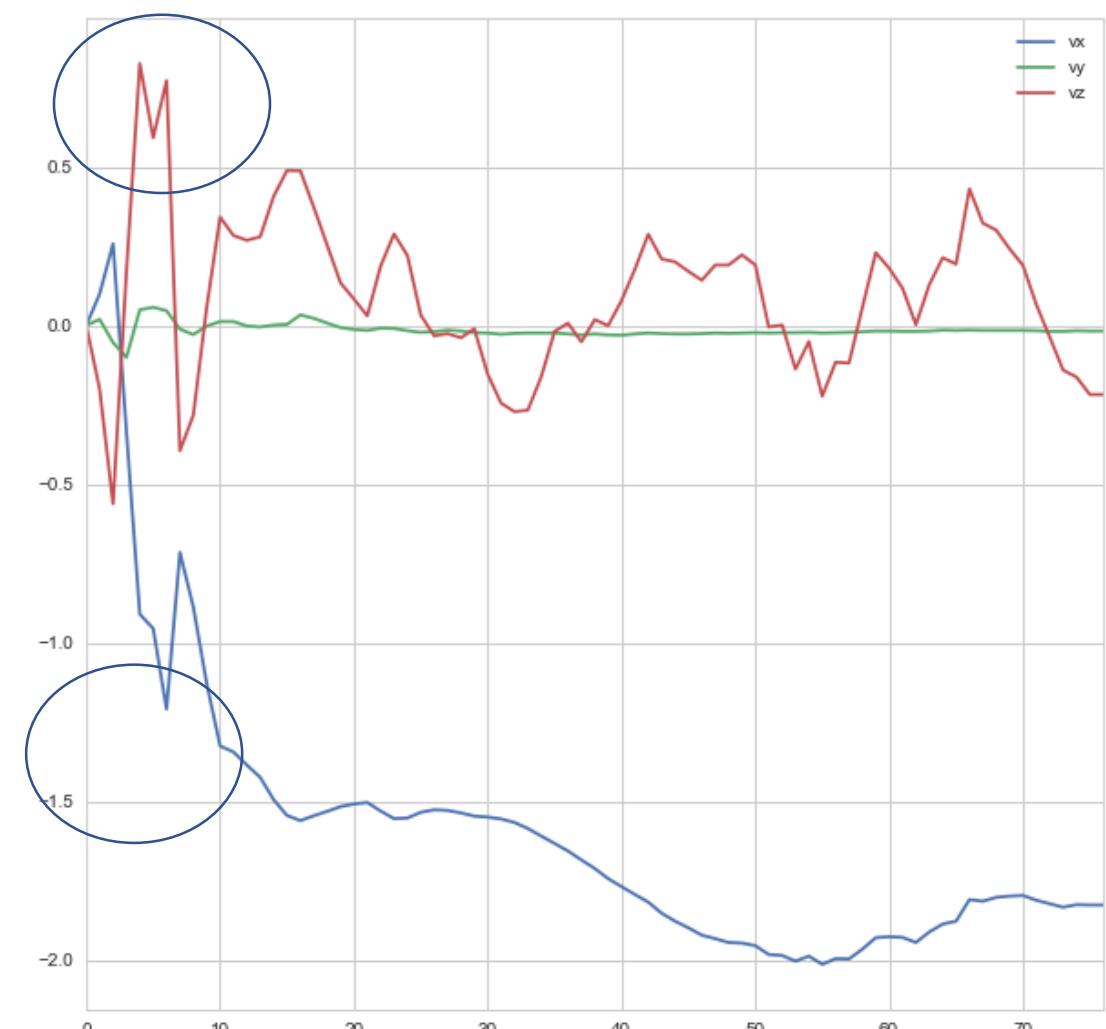
Body location / after



Body location / before



Body location / after



Research topic!

Note. The current filter architecture and parameters do their job, but certainly are not optimal. Field for further study!

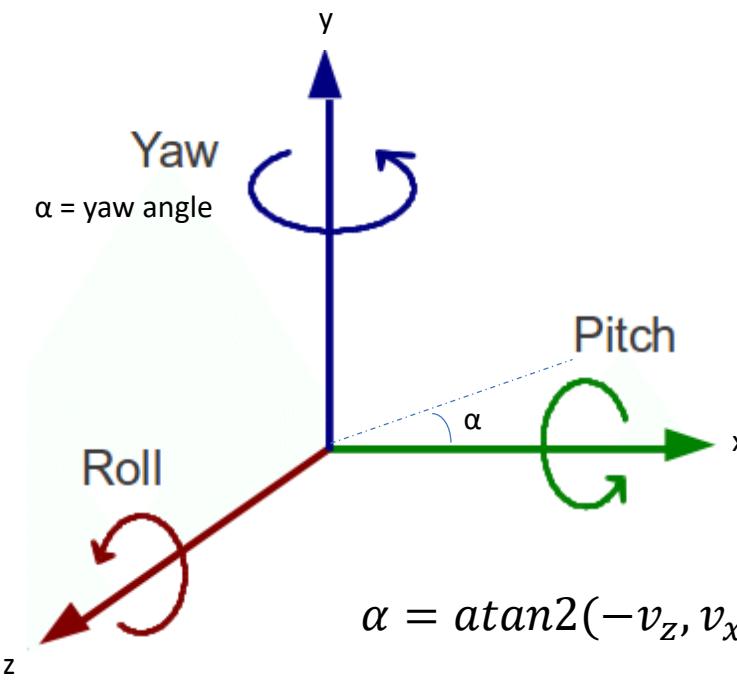


Body Volume and Orientation

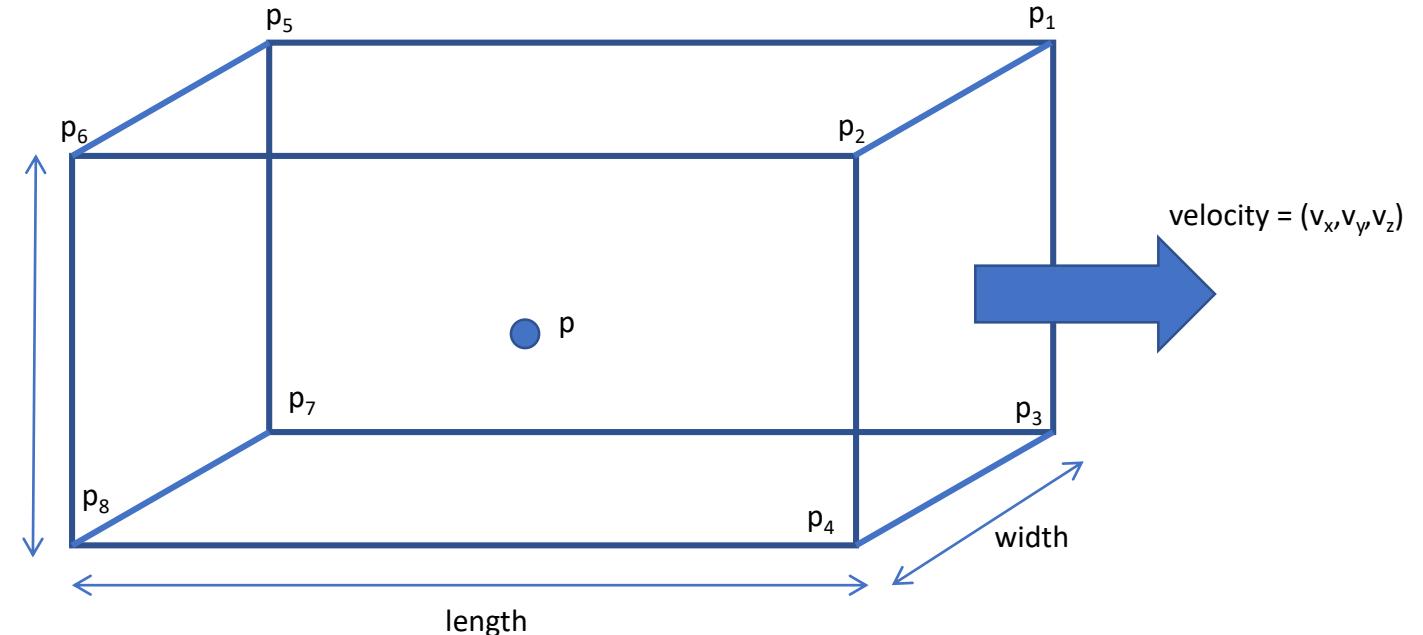
- Bounding sphere is not a good shape for volumes like car, human or train...
- KITTI uses bounded 3D box



Assumption: Only yaw allowed



Each body is supposed to have “normal” movement direction, which defines length, height and width.

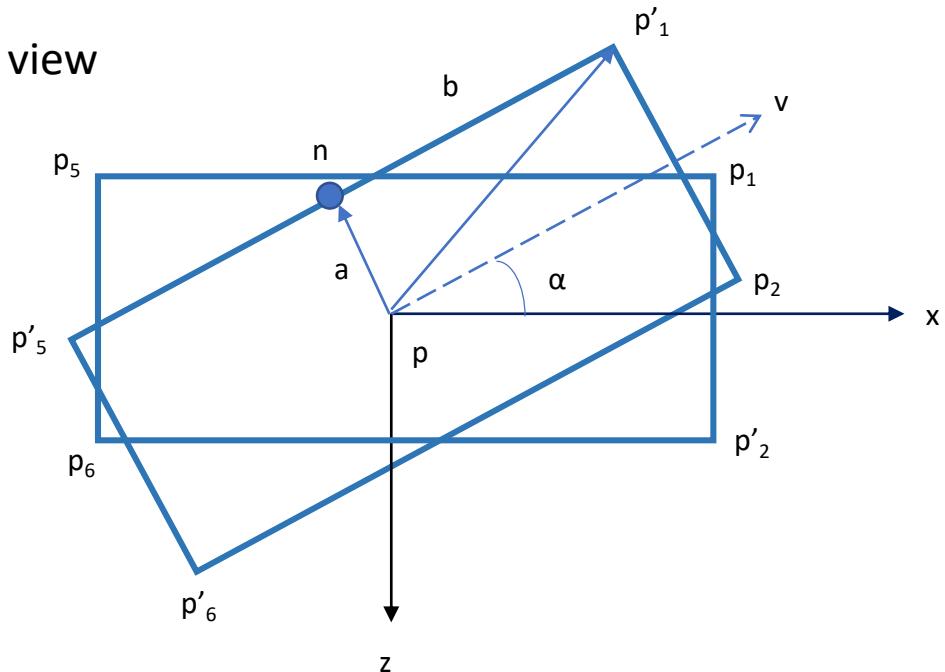


$$(p_{1,x}, p_{1,y}, p_{1,z}) = (p_x + 0.5 * \text{length}, p_y + 0.5 * \text{height}, p_z - 0.5 * \text{width})$$

...

$$(p_{8,x}, p_{8,y}, p_{8,z}) = (p_x - 0.5 * \text{length}, p_y - 0.5 * \text{height}, p_z + 0.5 * \text{width})$$

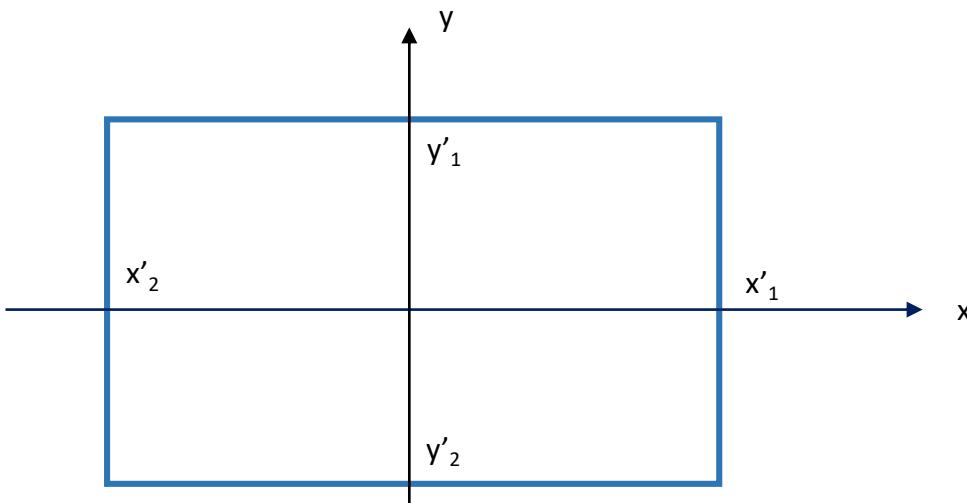
Top view



$$\alpha = \text{atan}2(-v_z, v_x)$$

$$\begin{bmatrix} p'_{x,i} \\ p'_{y,i} \\ p'_{z,i} \end{bmatrix} = \begin{bmatrix} \cos(\alpha) & 0 & \sin(\alpha) \\ 0 & 1 & 0 \\ -\sin(\alpha) & 0 & \cos(\alpha) \end{bmatrix} * \begin{bmatrix} p_{x,i} \\ p_{y,i} \\ p_{z,i} \end{bmatrix}$$

Side view

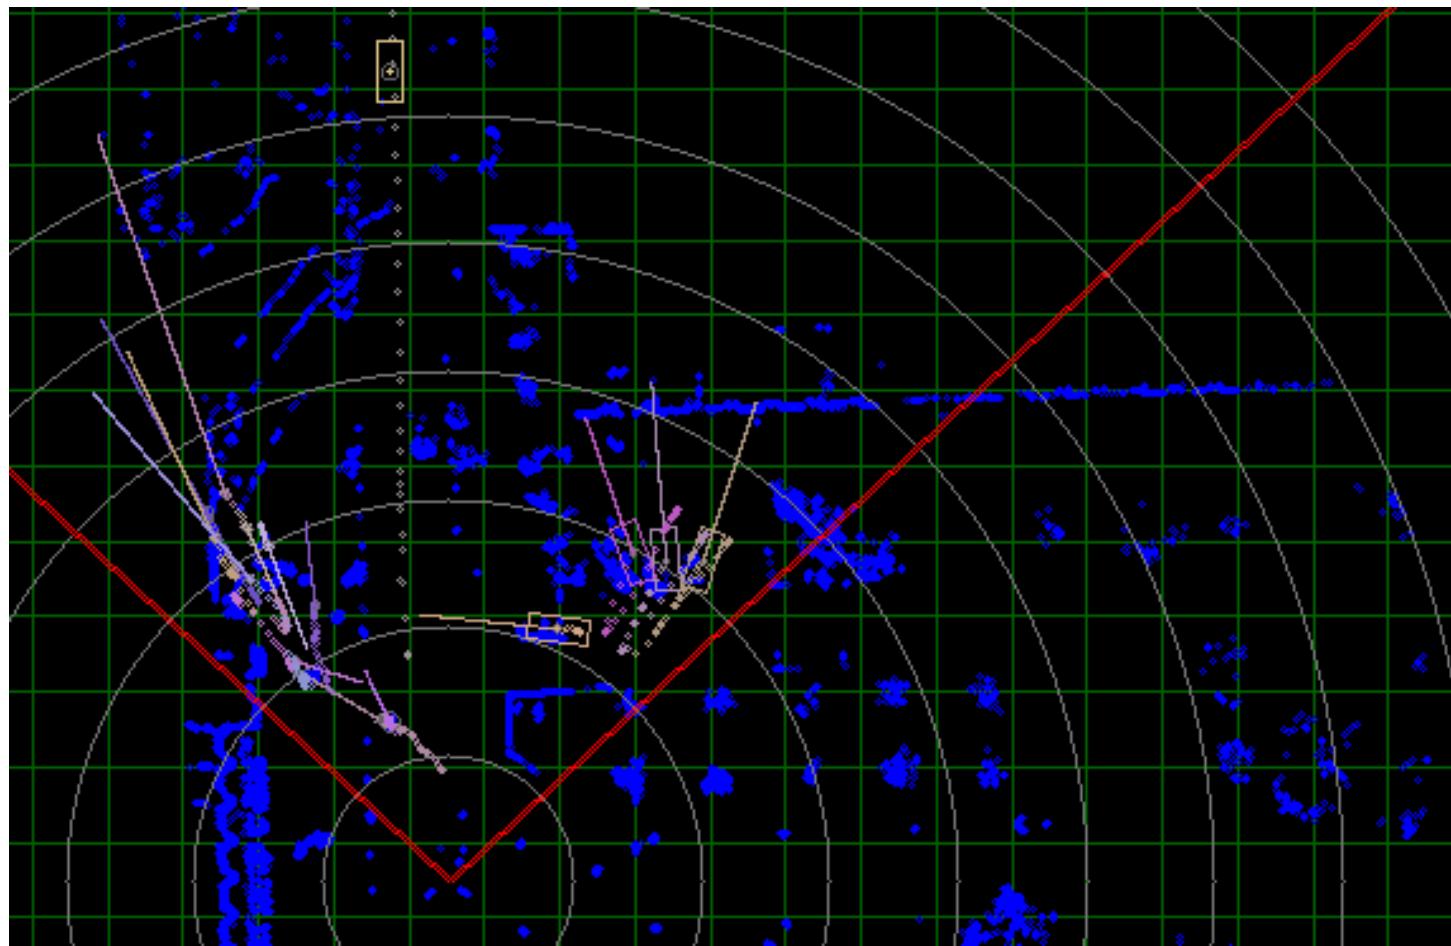


$$y'_1 = \text{height}/2$$

$$y'_2 = -\text{height}/2$$

$$x'_1 = \max(x_1, x_2, x_3, x_4, x_5, x_6, x_7, x_8)$$

$$x'_2 = \min(x_1, x_2, x_3, x_4, x_5, x_6, x_7, x_8)$$

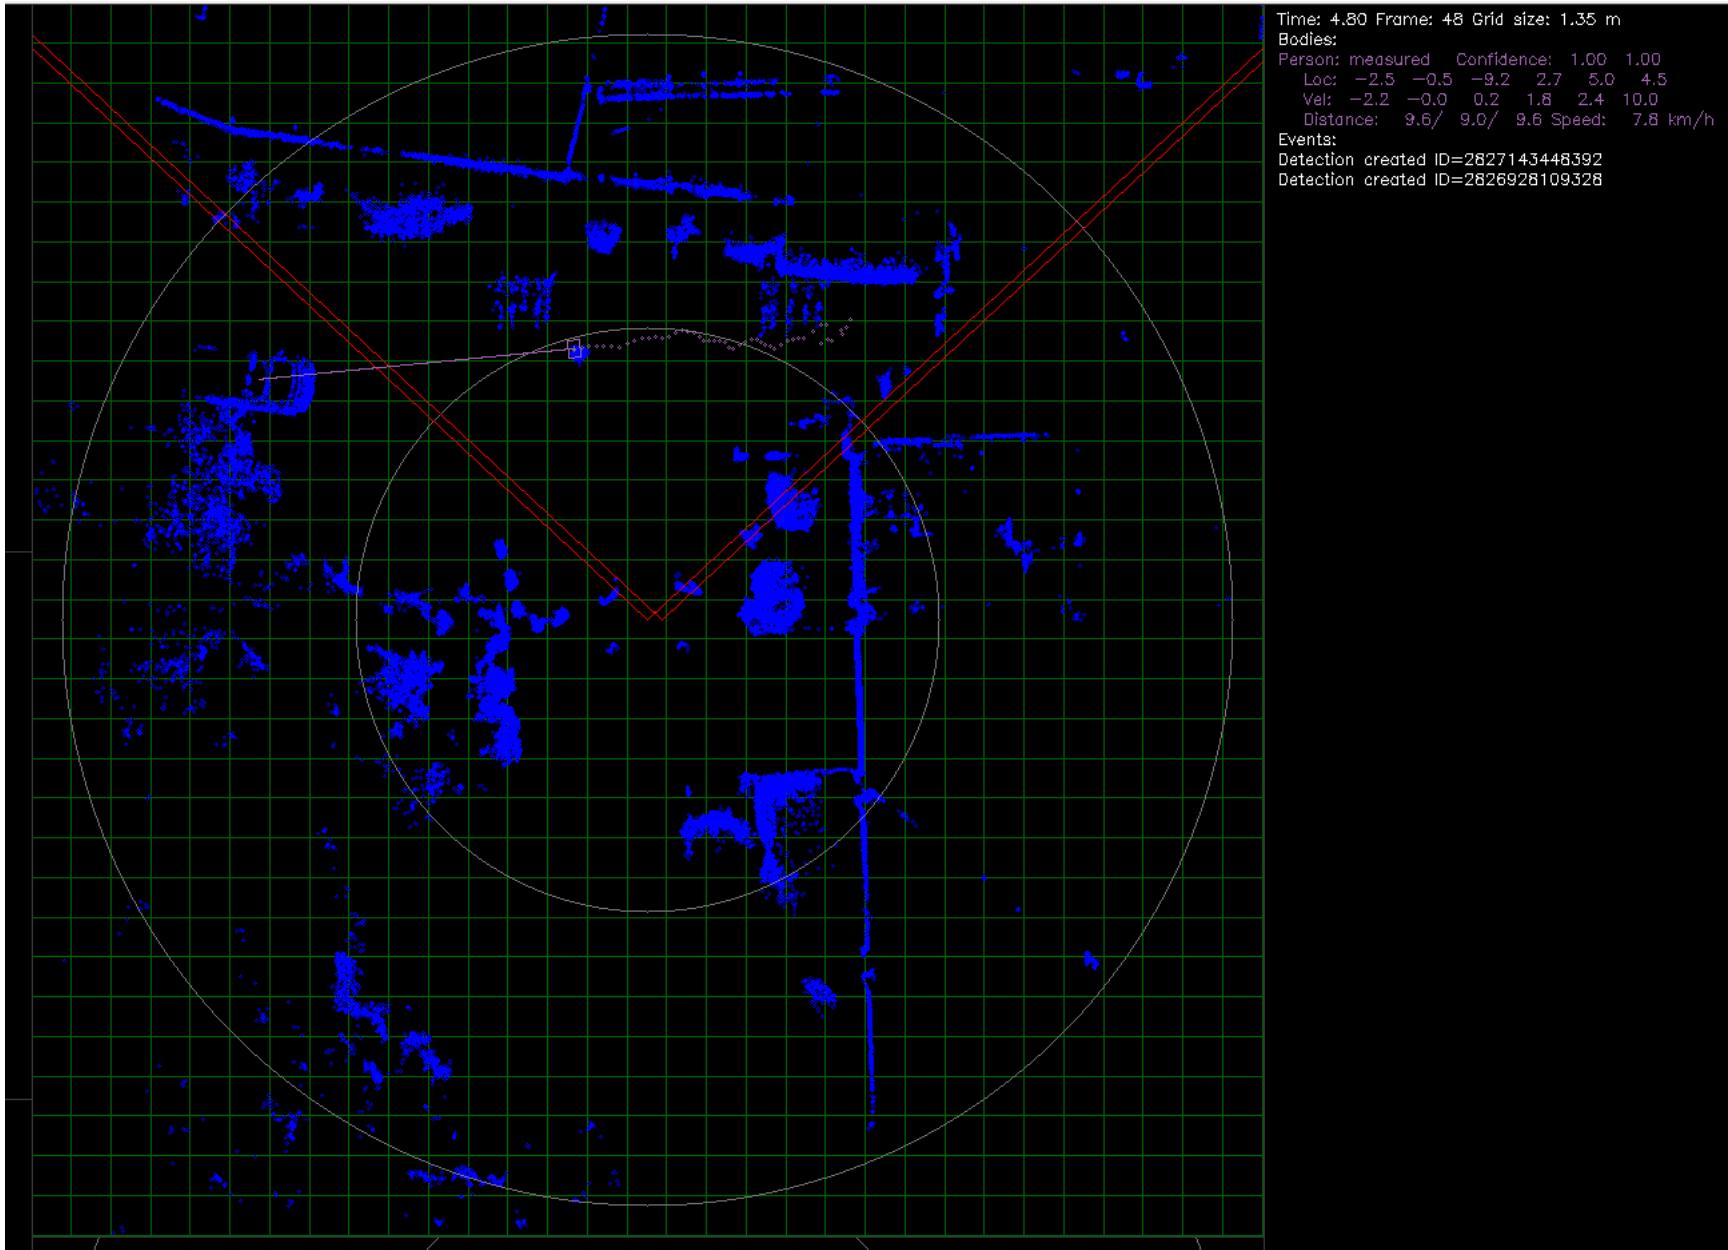


Currently used for display. Collision detection still uses spheres. TBD

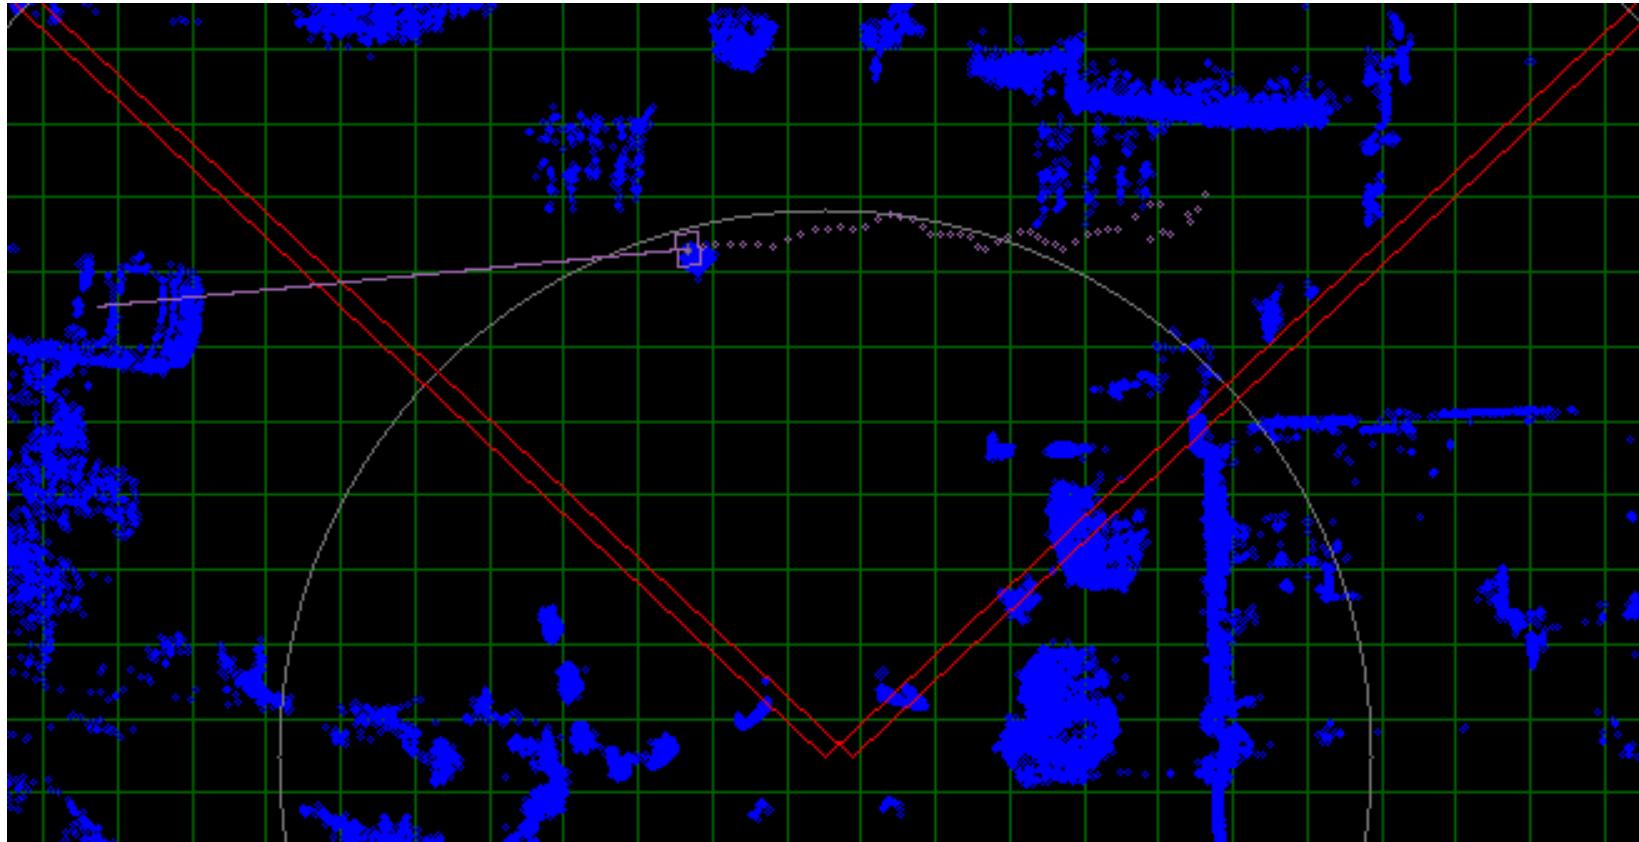
Grid Representation

Background:

Representing body movement forecast and probabilities in continuous world is presumably much more difficult than in discrete space. Grid is needed...



2D grid is used, with evenly spaced lines.
GRID_COUNT parameter (=31)



Representing conditional path statistics and probabilities is easier with grid.

Note. KITTI dataset includes "Person" category which is well suited to do research.

Data Category: Person

Before browsing, please wait some moments until this page is fully loaded.



2011_09_28_drive_0053 (0.3 GB)

Length: 74 frames (00:07 minutes)

Image resolution: 1392 x 512 pixels

Labels: 0 Cars, 0 Vans, 0 Trucks, 0 Pedestrians, 0 Sitters, 0 Cyclists, 0 Trams, 0 Misc
Downloads: [\[unsynced+unrectified data\]](#) [\[synced+rectified data\]](#) [\[calibration\]](#)



2011_09_28_drive_0054 (0.2 GB)

Length: 51 frames (00:05 minutes)

Image resolution: 1392 x 512 pixels

Labels: 0 Cars, 0 Vans, 0 Trucks, 0 Pedestrians, 0 Sitters, 0 Cyclists, 0 Trams, 0 Misc
Downloads: [\[unsynced+unrectified data\]](#) [\[synced+rectified data\]](#) [\[calibration\]](#)



2011_09_28_drive_0057 (0.3 GB)

Length: 80 frames (00:08 minutes)

Image resolution: 1392 x 512 pixels

Labels: 0 Cars, 0 Vans, 0 Trucks, 0 Pedestrians, 0 Sitters, 0 Cyclists, 0 Trams, 0 Misc
Downloads: [\[unsynced+unrectified data\]](#) [\[synced+rectified data\]](#) [\[calibration\]](#)



2011_09_28_drive_0065 (0.2 GB)

Length: 45 frames (00:04 minutes)

Image resolution: 1392 x 512 pixels

Labels: 0 Cars, 0 Vans, 0 Trucks, 0 Pedestrians, 0 Sitters, 0 Cyclists, 0 Trams, 0 Misc
Downloads: [\[unsynced+unrectified data\]](#) [\[synced+rectified data\]](#) [\[calibration\]](#)



2011_09_28_drive_0066 (0.1 GB)

Length: 35 frames (00:03 minutes)

Image resolution: 1392 x 512 pixels

Labels: 0 Cars, 0 Vans, 0 Trucks, 0 Pedestrians, 0 Sitters, 0 Cyclists, 0 Trams, 0 Misc
Downloads: [\[unsynced+unrectified data\]](#) [\[synced+rectified data\]](#) [\[calibration\]](#)



2011_09_28_drive_0068 (0.3 GB)

Length: 73 frames (00:07 minutes)

Image resolution: 1392 x 512 pixels

Labels: 0 Cars, 0 Vans, 0 Trucks, 0 Pedestrians, 0 Sitters, 0 Cyclists, 0 Trams, 0 Misc
Downloads: [\[unsynced+unrectified data\]](#) [\[synced+rectified data\]](#) [\[calibration\]](#)



2011_09_28_drive_0070 (0.2 GB)

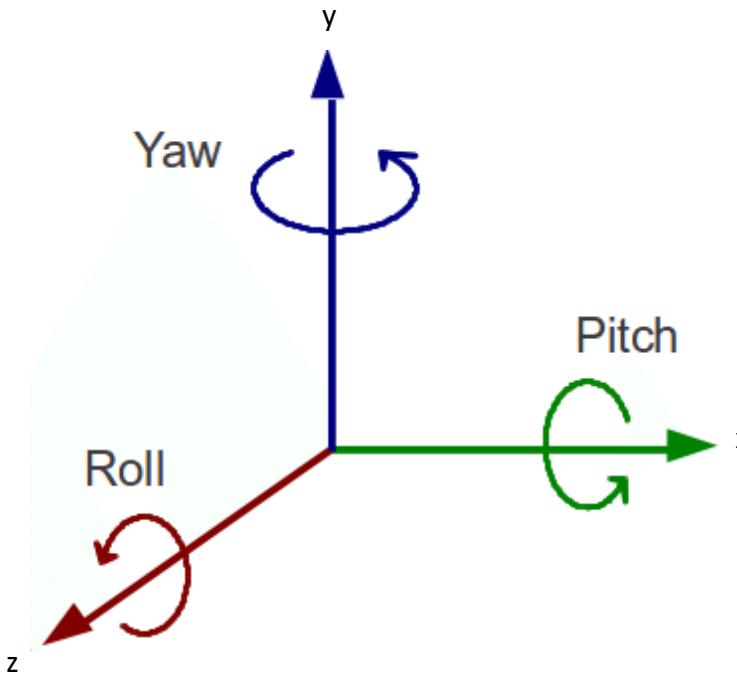
Length: 45 frames (00:04 minutes)



Research topic!

Nearly 100 videos where one or more persons are walking using different paths across the same area.

World Coordinates, Yaw Estimation



Assumptions:

- initially world coordinates coincide with left camera
- y, roll and pitch are fixed
- x and z are estimated from odometry (GPS)
- yaw is estimated from optical flow, or given

Camera has three degrees of freedom:

- x
- z
- yaw

Transformation from (left) camera coordinates into world coordinates:

Rotation + translation. Also known as 3D *rigid body motion* or the 3D *Euclidean transformation*, it can be written as $\mathbf{x}' = \mathbf{R}\mathbf{x} + \mathbf{t}$ or

$$\mathbf{x}' = \begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix} \bar{\mathbf{x}} \quad (2.24)$$

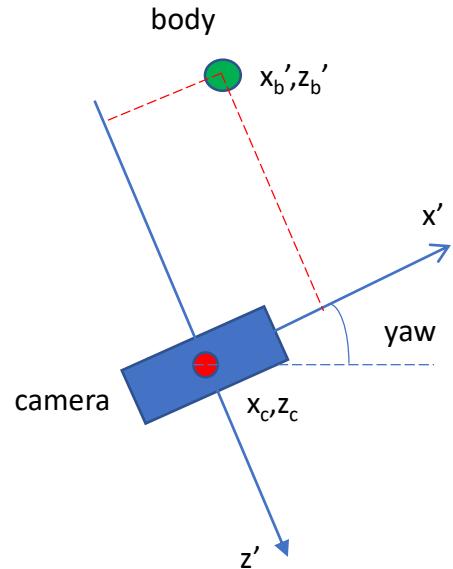
(Szeliski)

where \mathbf{R} is a 3×3 orthonormal rotation matrix with $\mathbf{R}\mathbf{R}^T = \mathbf{I}$ and $|\mathbf{R}| = 1$. Note that sometimes it is more convenient to describe a rigid motion using

$$\mathbf{x}' = \mathbf{R}(\mathbf{x} - \mathbf{c}) = \mathbf{R}\mathbf{x} - \mathbf{R}\mathbf{c}, \quad (2.25)$$

where \mathbf{c} is the center of rotation (often the camera center).

c from odometry
R from optical flow



From camera coordinates into world coordinates:

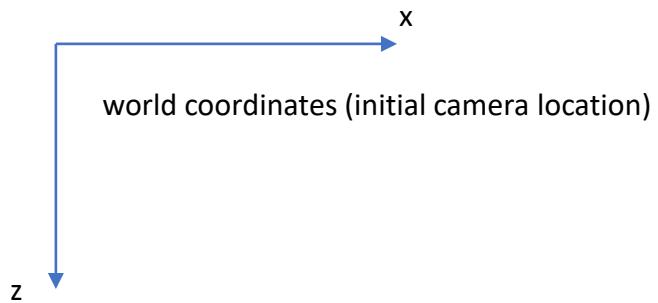
$$x_b = x_c + \cos(yaw) * x'_b + \sin(yaw) * z'_b$$

$$z_b = z_c - \sin(yaw) * x'_b + \cos(yaw) * z'_b$$

From world coordinates into camera coordinates:

$$x'_b = \cos(yaw) * (x_b - x_c) - \sin(yaw) * (z_b - z_c)$$

$$z'_b = \sin(yaw) * (x_b - x_c) + \cos(yaw) * (z_b - z_c)$$



If more degrees of freedom are wanted:

Which rotation representation is better?

The choice of representation for 3D rotations depends partly on the application.

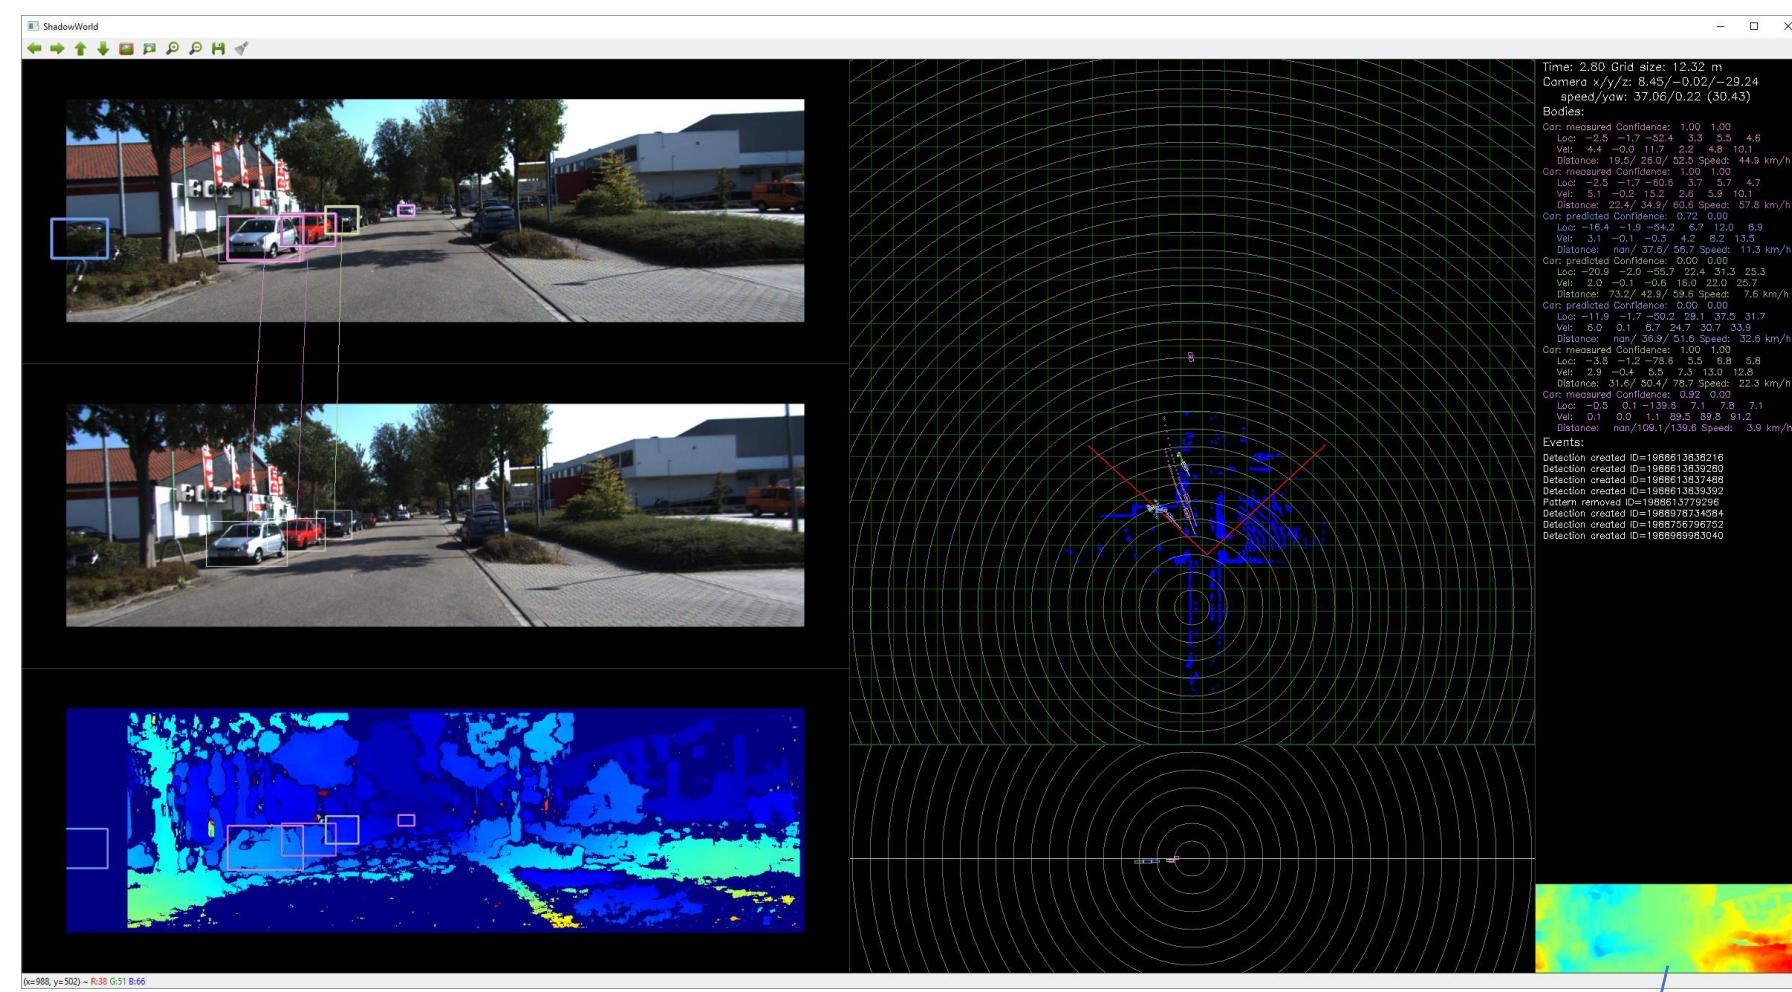
The axis/angle representation is minimal, and hence does not require any additional constraints on the parameters (no need to re-normalize after each update). If the angle is expressed in degrees, it is easier to understand the pose (say, 90° twist around x -axis), and also easier to express exact rotations. When the angle is in radians, the derivatives of \mathbf{R} with respect to ω can easily be computed (2.36).

Quaternions, on the other hand, are better if you want to keep track of a smoothly moving camera, since there are no discontinuities in the representation. It is also easier to interpolate between rotations and to chain rigid transformations (Murray, Li, and Sastry 1994; Bregler and Malik 1998).

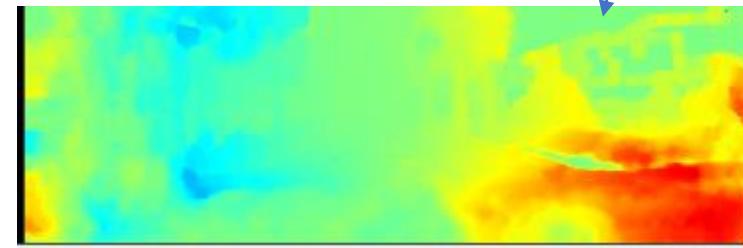
My usual preference is to use quaternions, but to update their estimates using an incremental rotation, as described in Section 6.2.2.

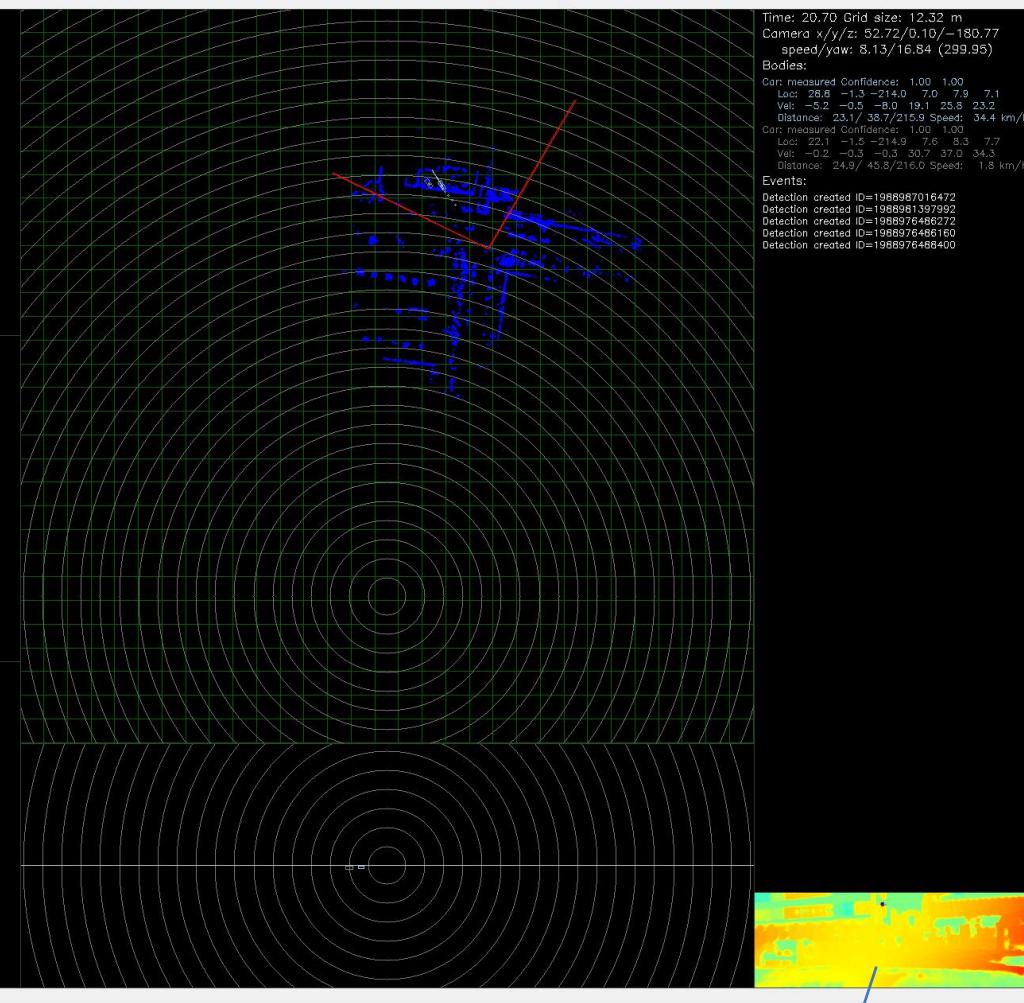
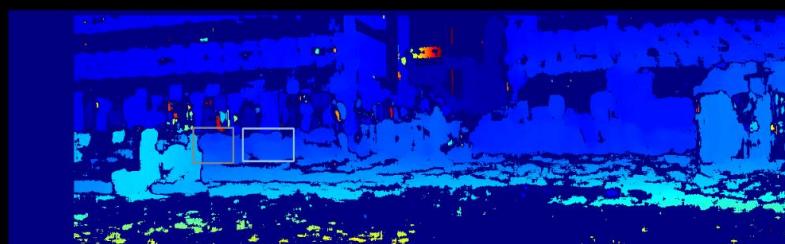
(Szeliski)

So, quaternions, but later when needed...

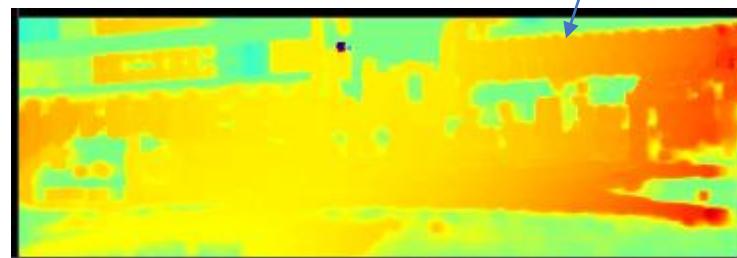


Car going straight ahead. Optical flow to left is coloured blue, and flow to right red. It is calculated using Farneback algorithm.

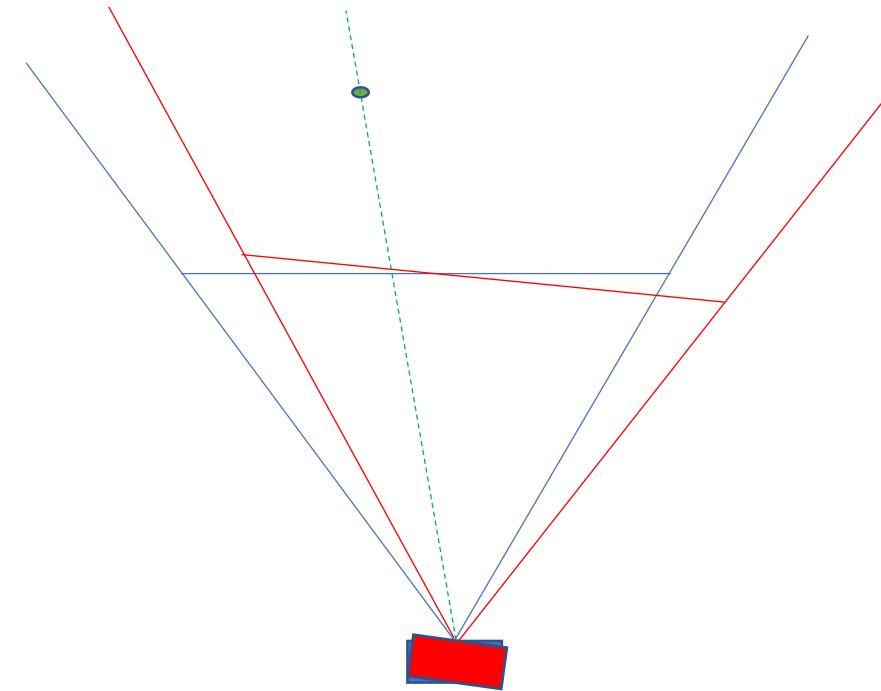




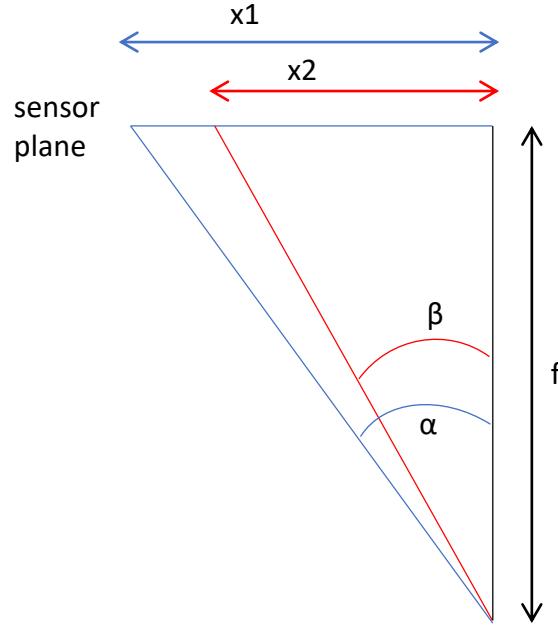
Car turning left. Optical flow is mainly to right and colored red/yellow.



From optical flow to yaw



Approximation for small turns:



$$x_1 = f * \tan(\alpha) \quad x_2 = f * \tan(\beta)$$

$$x_2 - x_1 = f * \tan(\alpha) - f * \tan(\beta)$$

$$\beta = \arctan(\tan(\alpha) - \frac{x_1 - x_2}{f})$$

Substituting $\alpha=0$ (reference):

$$\gamma = (\text{amount of turning}) = \beta - \alpha = \arctan\left(\frac{x_2 - x_1}{f}\right)$$

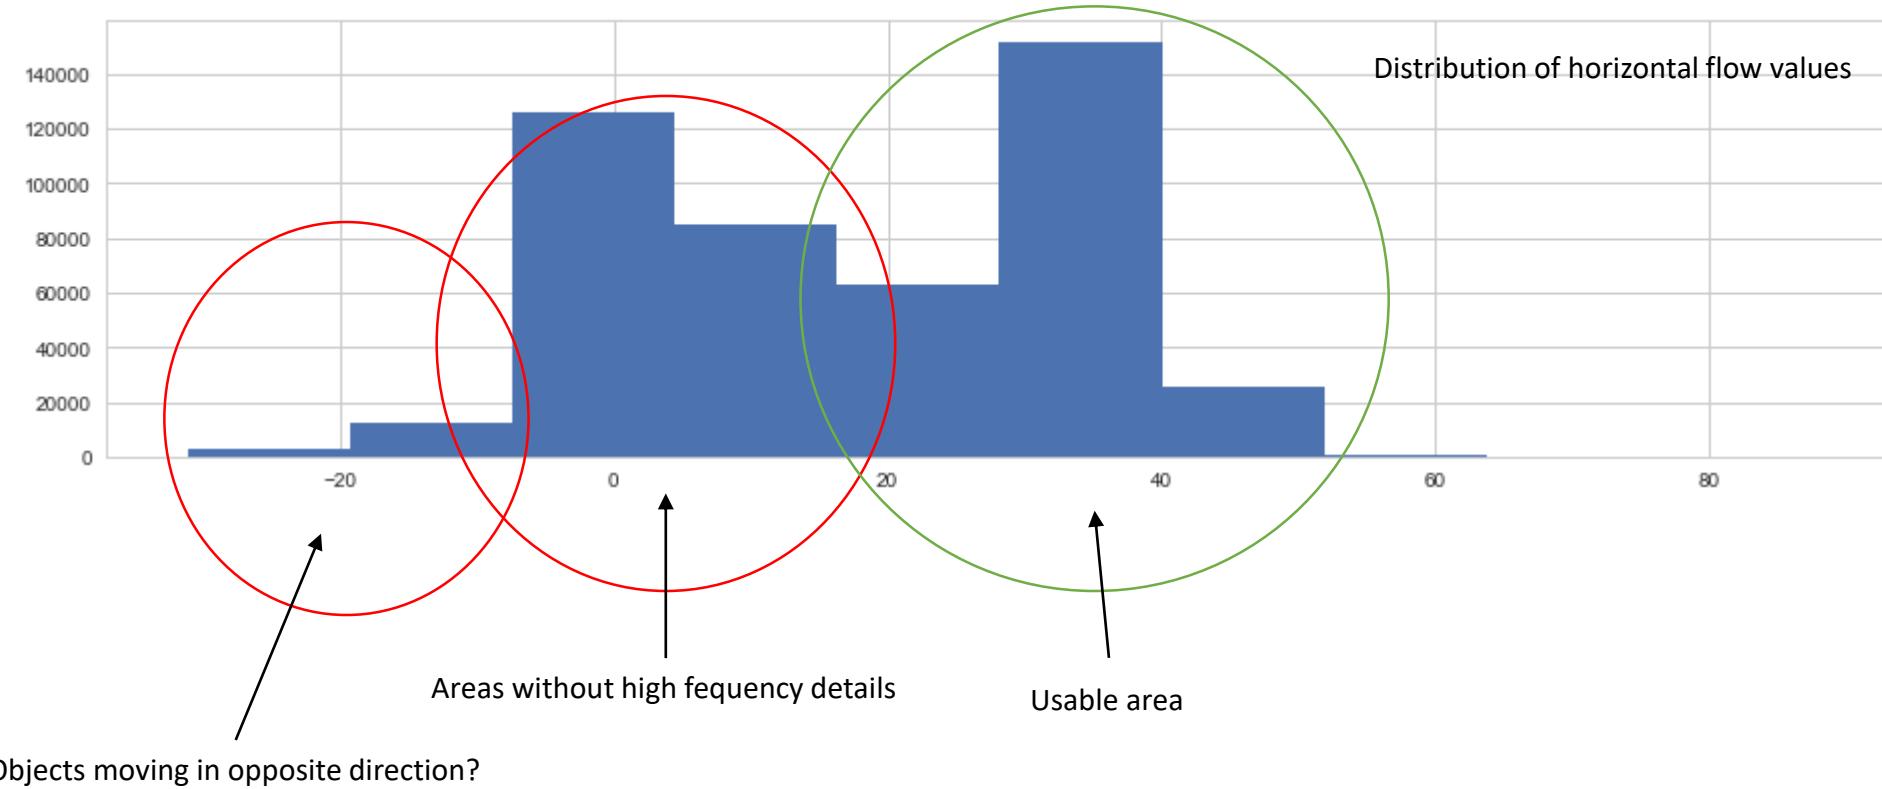
$$\gamma = (\text{amount of turning}) = \beta - \alpha = \arctan\left(\frac{x_2 - x_1}{f}\right) = \arctan\left(\frac{dx_p * sw}{f * pw}\right)$$

s_w = sensor width (m)
 p_w = image width (pixels)
 f = focal length (m)
 dx_p = horizontal optical flow in pixels

Cumulative yaw is maintained for the reference (left) camera:

$$yaw_n = \sum_{i=1}^n \gamma_i$$

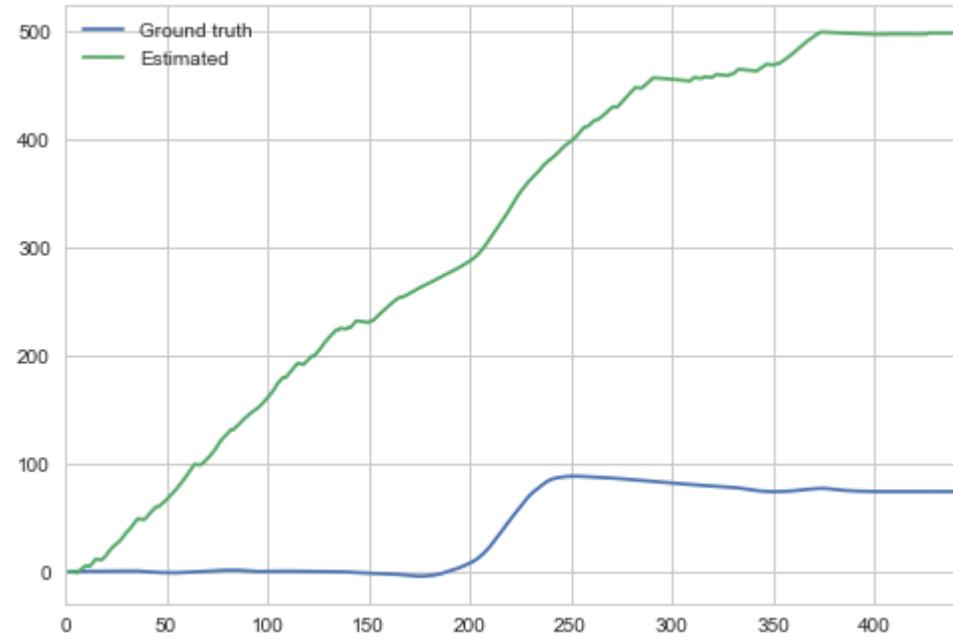
Note: Approximation and optical flow errors are cumulative over time. However, the major effect of error is that the initial reference coordinate system is turning in space around y-axis. As this is completely arbitrary, it doesn't matter much. If we would like to map environment, it would matter more.



Objects moving in opposite direction?

Ad hoc algorithm to determine 'average' background optical flow:

1. Determine main direction (+/-) by calculating mean of flows
2. Drop all flows in the opposite direction
3. Calculate mean value of flows in right direction
4. Drop all flows whose absolute value is below mean value
5. Determine the final value by calculating mean value of remaining flows



Naive ad hoc algorithm is not very good, but a starting point...

Note 1. Lucas-Kanade algorithm is another possibility. It concentrates on areas with texture and high frequency details. Anyway, areas corresponding to patterns should be eliminated.

Note 2. KITTI dataset includes 6 DOF data for moving camera. Lots of data to do research!!!

```
dataformat.txt
1 lat:   latitude of the oxts-unit (deg)
2 lon:   longitude of the oxts-unit (deg)
3 alt:   altitude of the oxts-unit (m)
4 roll:  roll angle (rad),    0 = level, positive = left side up,      range: -pi .. +pi
5 pitch: pitch angle (rad),  0 = level, positive = front down,       range: -pi/2 .. +pi/2
6 yaw:   heading (rad),     0 = east,  positive = counter clockwise, range: -pi .. +pi
7 vn:    velocity towards north (m/s)
8 ve:    velocity towards east (m/s)
9 vf:    forward velocity, i.e. parallel to earth-surface (m/s)
10 vl:   leftward velocity, i.e. parallel to earth-surface (m/s)
11 vu:   upward velocity, i.e. perpendicular to earth-surface (m/s)
12 ax:   acceleration in x, i.e. in direction of vehicle front (m/s^2)
13 ay:   acceleration in y, i.e. in direction of vehicle left (m/s^2)
14 az:   acceleration in z, i.e. in direction of vehicle top (m/s^2)
15 af:   forward acceleration (m/s^2)
16 al:   leftward acceleration (m/s^2)
17 au:   upward acceleration (m/s^2)
18 wx:   angular rate around x (rad/s)
19 wy:   angular rate around y (rad/s)
20 wz:   angular rate around z (rad/s)
21 wf:   angular rate around forward axis (rad/s)
22 wl:   angular rate around leftward axis (rad/s)
23 wu:   angular rate around upward axis (rad/s)
24 pos_accuracy: velocity accuracy (north/east in m)
25 vel_accuracy: velocity accuracy (north/east in m/s)
26 navstat:   navigation status (see navstat_to_string)
27 numstsats: number of satellites tracked by primary GPS receiver
28 posmode:   position mode of primary GPS receiver (see gps_mode_to_string)
29 velmode:   velocity mode of primary GPS receiver (see gps_mode_to_string)
30 orimode:   orientation mode of primary GPS receiver (see gps_mode_to_string)
```



Research topic!

COCO Dataset

Common Objects in Context



News

- 2017 Challenge Winners for Detection, Keypoint, & Stuff tasks have been announced! Please visit the [Joint COCO and Places Recognition ICCV workshop page](#) for details.
- This website is now hosted on [Github](#), which provides page source and history.
- Keypoint analysis tools are now available, see [Keypoints evaluation](#), Section 4.

What is COCO?



COCO is a large-scale object detection, segmentation, and captioning dataset. COCO has several features:

- ✓ Object segmentation
- ✓ Recognition in context
- ✓ Superpixel stuff segmentation
- ✓ 330K images (>200K labeled)
- ✓ 1.5 million object instances
- ✓ 80 object categories
- ✓ 91 stuff categories
- ✓ 5 captions per image
- ✓ 250,000 people with keypoints

Collaborators

Tsung-Yi Lin Google Brain
Genevieve Patterson MSR
Matteo R. Ronchi Caltech
Yin Cui Cornell Tech
Michael Maire TTI-Chicago
Serge Belongie Cornell Tech
Lubomir Bourdev WaveOne, Inc.
Ross Girshick FAIR
James Hays Georgia Tech
Pietro Perona Caltech
Deva Ramanan CMU
Larry Zitnick FAIR
Piotr Dollár FAIR

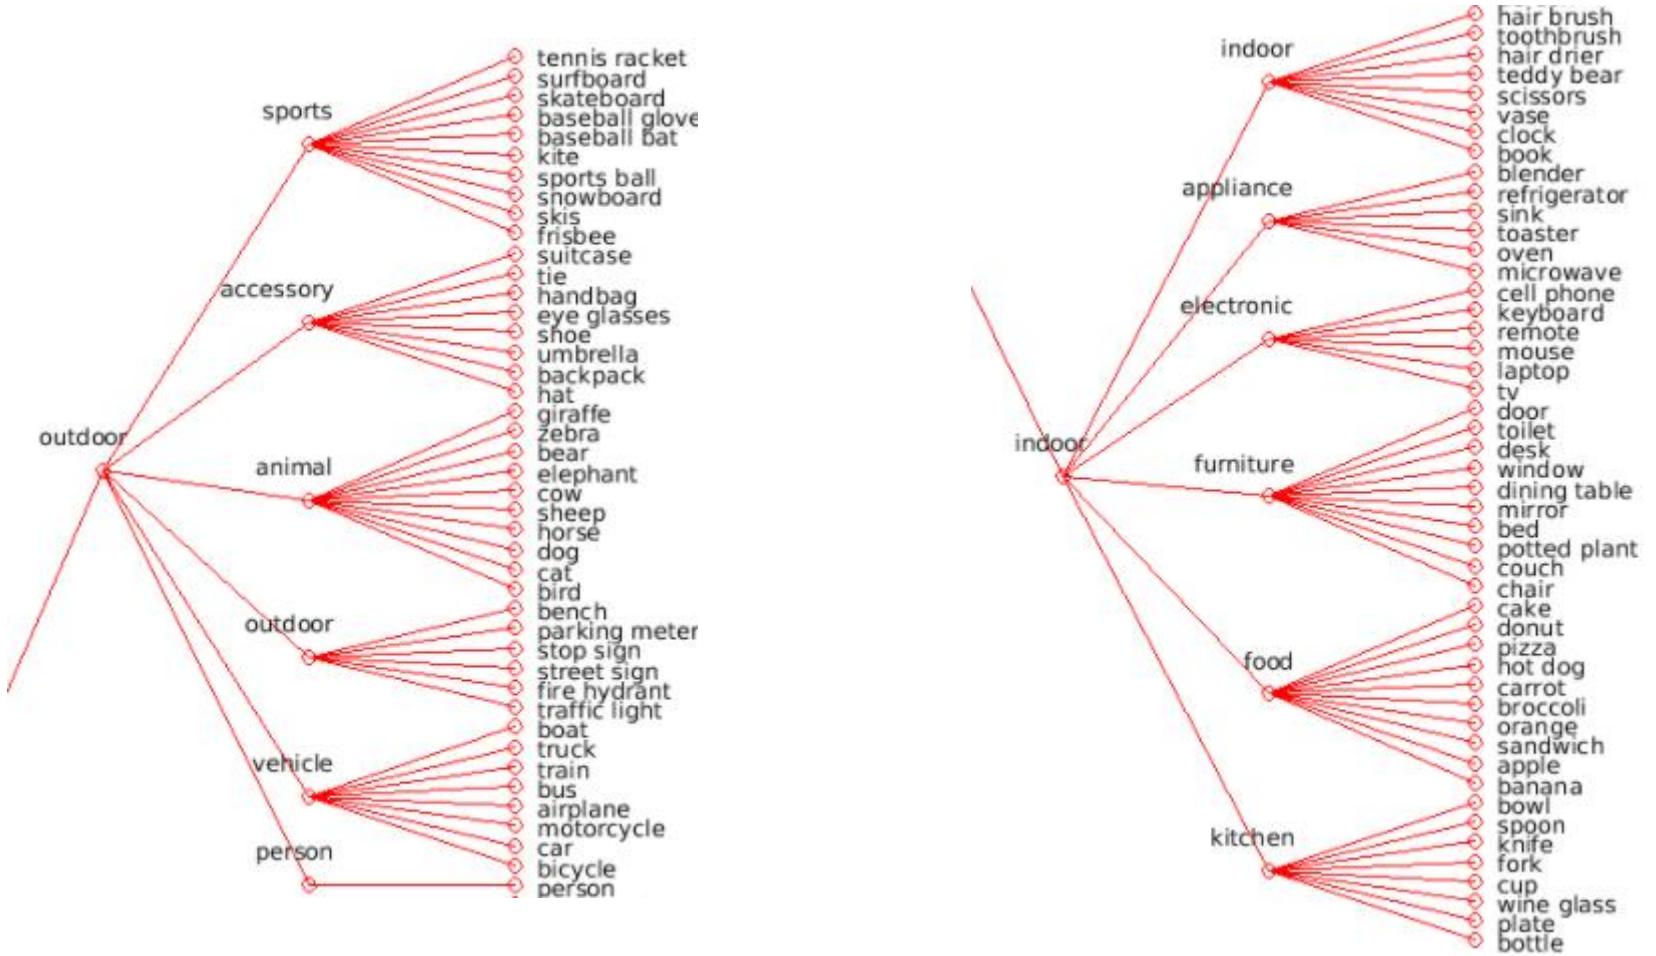
Sponsors



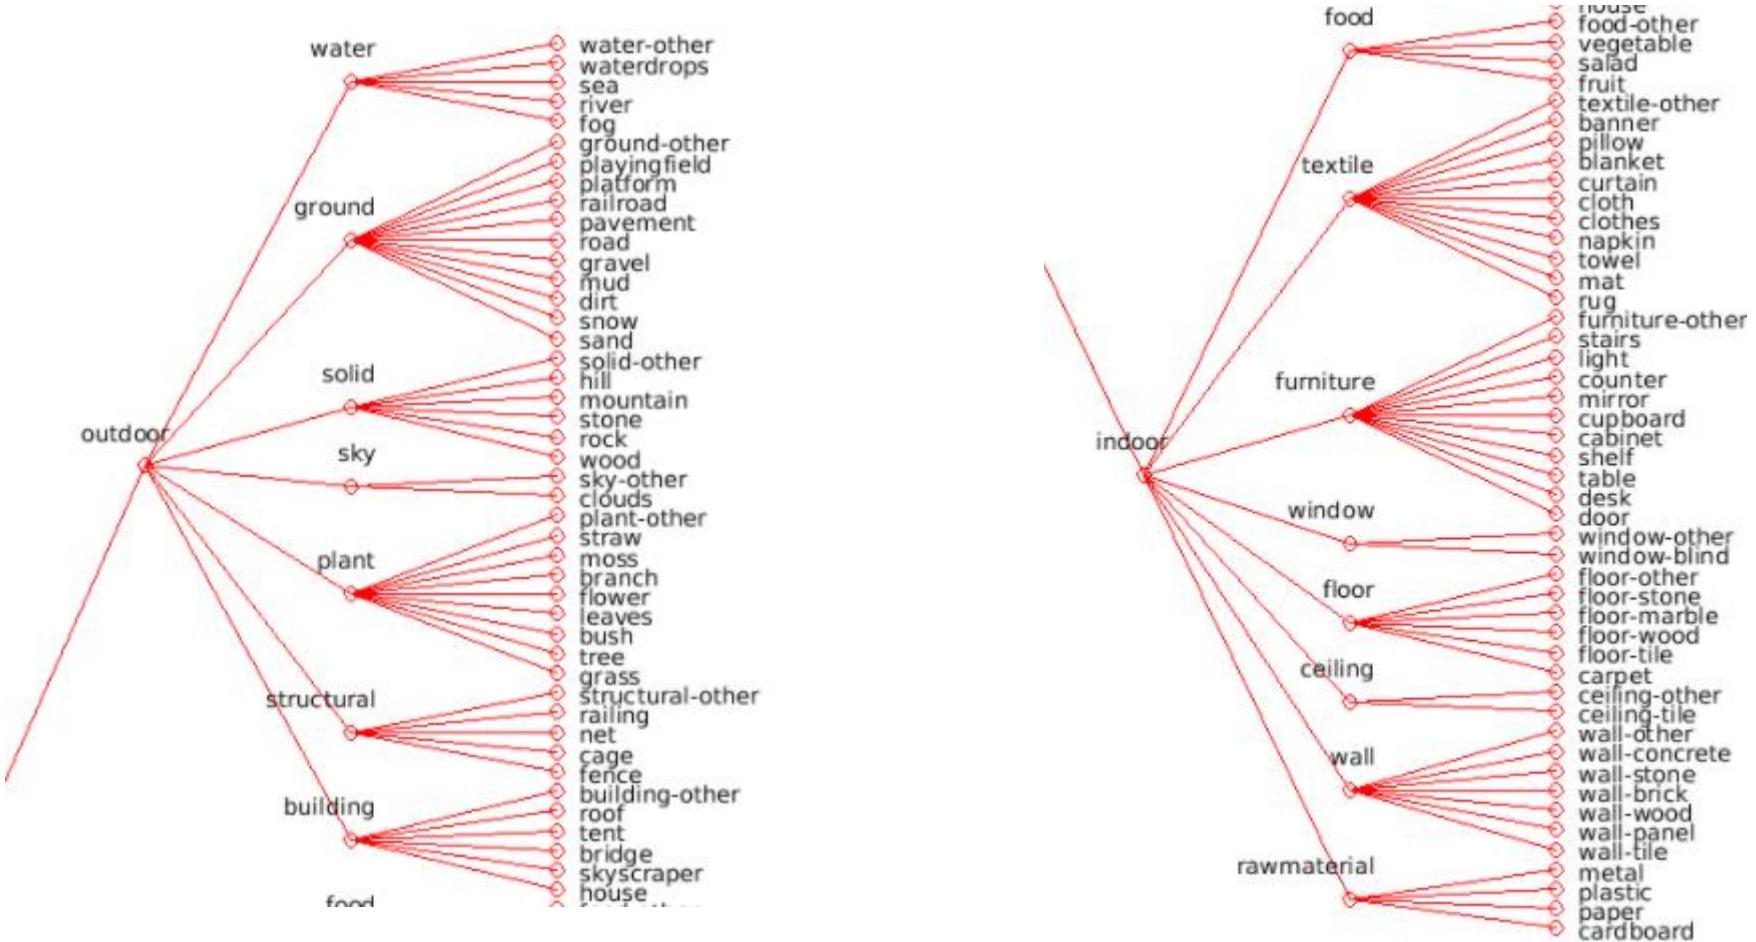
Why?

- 171 everyday common categories
- Large amount of images
- Good API
- Emerging support for implementations
- Captions

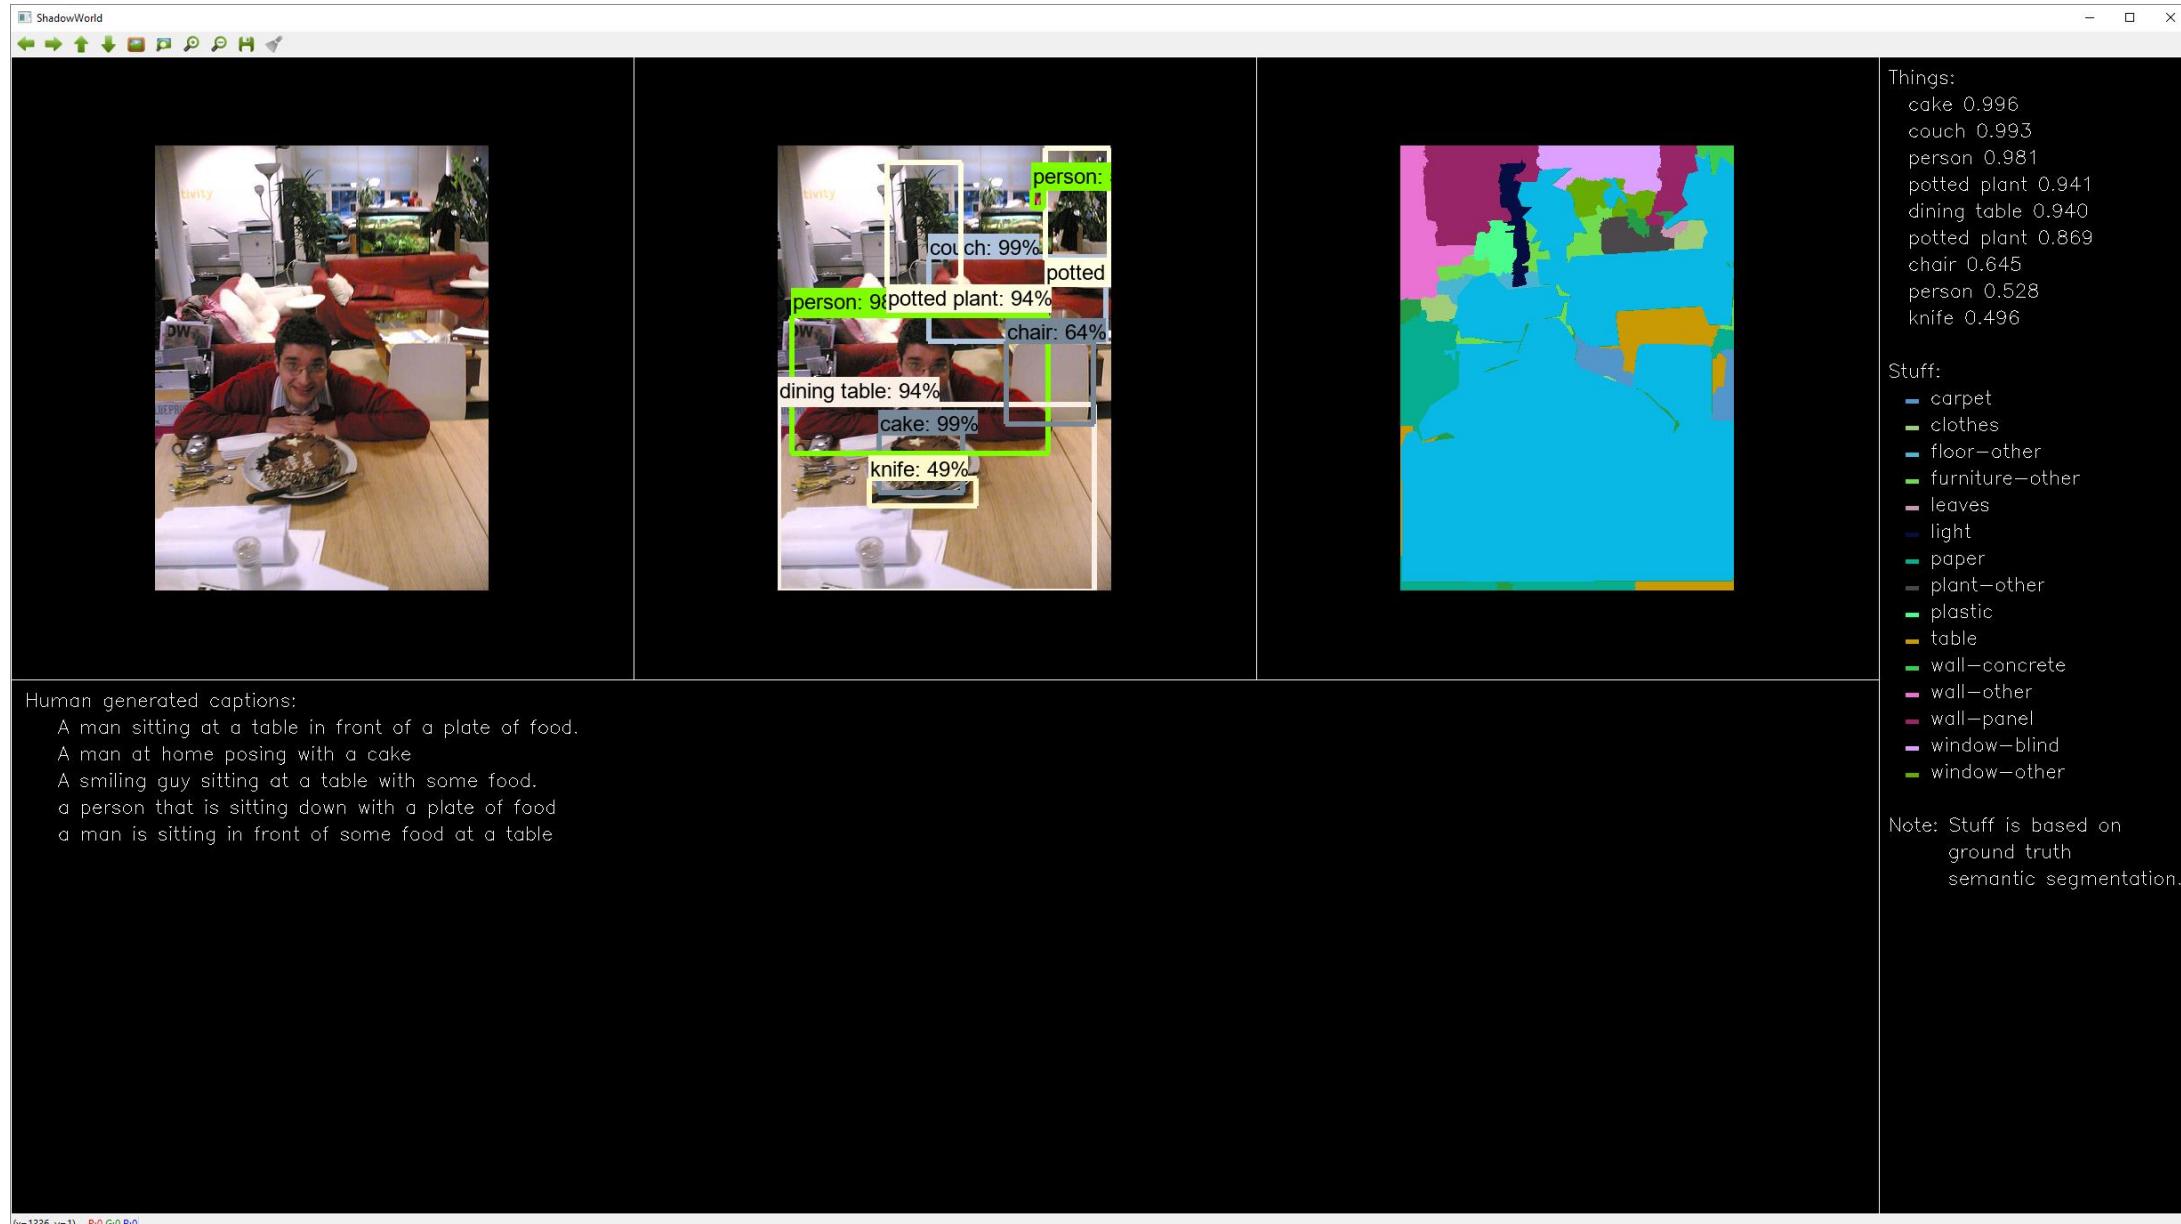
COCO things (=bodies)




COCO stuff



ShadowWorld 4.0 user interface for COCO (with faster rcnn/resnet101 object detection network):



Semantic Segmentation



A large long **train** on a steel **track**.

A blue and yellow transit **train** leaving the **station**.

A **train** crossing beneath a city **bridge** with brick **towers**.

A **train** passing by an over **bridge** with a railway **track** (...).

A **train** is getting ready to leave the train **station**.

Why?

The goal is to give rich view description. For that, we need the bodies (things) and context (stuff).

Unfortunately, there is no TensorFlow implementation for COCO stuff detection. However, there are Caffe implementations. For those implementations, necessary files for converting them into OpenCV DNN could not be found.

Options to continue:

1. Wait until TensorFlow implementation is available.
2. Learn Caffe and export an implementation into TensorFlow (or construct a bridge from Caffe to TensorFlow).
3. Take existing TensorFlow implementation (for example DeepLab/Pascal VOC 2012) and use transfer learning to adapt into COCO.
4. Teach a TensorFlow network from scratch.

I chose to wait until July 2018. If there is no TensorFlow implementation at that time, option 2 is followed. If that fails, option 3 is taken. Option 4 would be very tempting, with a lot to learn. However, there are so many moving parts and hyperparameter adjustments in the process, that this time option 4 should be avoided. If it were followed, it would be a great research topic. It would also require using cloud services, which are not free. Data is available.



Until having a semantic segmentation implementation, the COCO ground truth will be used to develop natural language models.



To summarize:

Things are created by object detection, stuff using semantic segmentation.

Speech Recognition

SpeechRecognition 3.8.1 Python package
+
PocketSphinx

Google API Client Library for Python +
Google Cloud Speech API much more better, but not free.
Will be used in the prototype. Requires network connection.

Use Cases

[218v2] 29 Mar 2018

VizWiz Grand Challenge: Answering Visual Questions from Blind People

Danna Gurari¹, Qing Li², Abigale J. Stangl³, Anhong Guo⁴, Chi Lin¹,
Kristen Grauman¹, Jiebo Luo⁵, and Jeffrey P. Bigham⁴

¹ University of Texas at Austin, ² University of Science and Technology of China,
³ University of Colorado Boulder, ⁴ Carnegie Mellon University ⁵ University of Rochester


Abstract

The study of algorithms to automatically answer visual questions currently is motivated by visual question answering (VQA) datasets constructed in artificial VQA settings. We propose VizWiz, the first goal-oriented VQA dataset arising from a natural VQA setting. VizWiz consists of over 31,000 visual questions originating from blind people who each took a picture using a mobile phone and recorded a spoken question about it, together with 10 crowdsourced answers per visual question. VizWiz differs from the many existing VQA datasets because (1) images are captured by blind photographers and so are often poor quality, (2) questions are spoken and so are more conversational, and (3) often visual questions cannot be answered. Evaluation of modern algorithms for answering visual questions and deciding if a visual question is answerable reveals that VizWiz is a challenging dataset. We introduce this dataset to en-

question about any image.

Over the past three years, 14 VQA datasets have emerged in the vision community to catalyze research on the VQA problem [6, 7, 16, 17, 19, 20, 23, 28, 32, 39, 40, 43, 45]. Historically, progress in the research community on a given computer vision problem is typically preceded by a large-scale, publicly-shared dataset [12, 25, 30, 33, 42]. However, a limitation of available VQA datasets is that all come from artificially created VQA settings. Moreover, none are “goal oriented” towards the images and questions that come from blind people. Yet, blind people arguably have been producing the big data desired to train algorithms. For nearly a decade, blind people have been both taking pictures [4, 8] and asking questions about the pictures they take [8, 11, 24]. Moreover, blind people often are early adopters of computer vision tools to support their *real* daily needs.

We introduce the first publicly-available vision dataset originating from blind people, which we call “VizWiz”,



VizWiz Tutorial

Submitting a question using the VizWiz app



▶ ▶| 🔍 0:02 / 3:22

CC  □ ⌂

VizWiz Tutorial

11,109 views

 8 1 SHARE ...

OCR



Q: Does this foundation have any sunscreen?
A: yes



Q: What is this?
A: 10 euros

Attribute extraction



Q: What color is this?
A: green



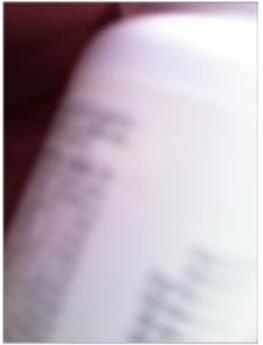
Q: Please can you tell me what this item is?
A: butternut squash red pepper soup



Q: Is it sunny outside?
A: yes



Q: Is this air conditioner on fan, dehumidifier, or air conditioning?
A: air conditioning



Q: What type of pills are these?
A: unsuitable image



Q: What type of soup is this?
A: unsuitable image

Requires specific network



Research topic!



Q: When is the expiration date?
A: unanswerable



Q: What is this?
A: unanswerable



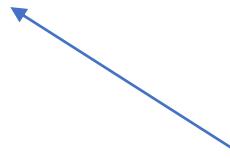
Q: Can you please tell me what the oven temperature is set to?
A: unanswerable

Impossible?




Figure 2. Distribution of the first six words for all questions in VizWiz. The innermost ring represents the first word and each subsequent ring represents a subsequent word. The arc size is proportional to the number of questions with that initial word/phrase.

Note: Because VizWiz is based on q/a data on specific mobile app, some questions are absent, probably because people know it is unable to give the answer. Most notably, questions beginning "where". This is where image based situation awareness shines! But it also means that we cannot use pure VizWiz as learning and evaluation dataset. Another set of questions having only small number of representatives is "How many...?"



The goal will be to cover large amount of this this circle + some extra questions.

ShadowWorld integration (easy example, more behind the link on Youtube video...)



Dataset Roles

KITTI

- 3D object localization
- State estimation
- Forecasting

COCO


- Object and stuff classes
- Network implementations

VIZWIZ

- Use cases

New Ideas

Using masks instead of (or with) bounding boxes:



Mask can be applied to conventional stereo disparity map to get many distance measurements. Object attributes more accurate.

TensorFlow Object Detection Model Zoo

COCO-trained models {#coco-models}

Model name	Speed (ms)	COCO mAP[^1]	Outputs
ssd_mobilenet_v1_coco	30	21	Boxes
ssd_mobilenet_v2_coco	31	22	Boxes
ssd_inception_v2_coco	42	24	Boxes
faster_rcnn_inception_v2_coco	58	28	Boxes
faster_rcnn_resnet50_coco	89	30	Boxes
faster_rcnn_resnet50_lowproposals_coco	64		Boxes
rfcn_resnet101_coco	92	30	Boxes
faster_rcnn_resnet101_coco	106	32	Boxes
faster_rcnn_resnet101_lowproposals_coco	82		Boxes
faster_rcnn_inception_resnet_v2_atrous_coco	620	37	Boxes
faster_rcnn_inception_resnet_v2_atrous_lowproposals_coco	241		Boxes
faster_rcnn_nas	1833	43	Boxes
faster_rcnn_nas_lowproposals_coco	540		Boxes
mask_rcnn_inception_resnet_v2_atrous_coco	771	36	Masks
mask_rcnn_inception_v2_coco	79	25	Masks
mask_rcnn_resnet101_atrous_coco	470	33	Masks
mask_rcnn_resnet50_atrous_coco	343	29	Masks

Current network

Mask networks

Particle cloud representation for both things (bodies) and stuff

- Volumes
- Helps probabilistic reasoning

Grid (2D or 3D) world representation

- Enables set of methods and tools developed for grid representation

Next Steps

Discussion

Thank you!

lampola@student.tut.fi

<https://github.com/SakariLampola/Thesis>